

# Circuit Codex

## User manual

What each tool does, what its terms mean, and where its numbers come from. Tap a tool to go to its section.

# Index

## PASSIVE COMPONENTS

|                         |            |
|-------------------------|------------|
| Resistor color code     | Resistors  |
| SMD resistor code       | Resistors  |
| Resistor power rating   | Resistors  |
| E-series values         | Resistors  |
| Series/parallel (R)     | Resistors  |
| Voltage divider         | Resistors  |
| Current divider         | Resistors  |
| Wheatstone bridge       | Resistors  |
| Delta-Y transform       | Resistors  |
| NTC/PTC thermistor      | Resistors  |
| Ceramic cap code        | Capacitors |
| Film capacitor code     | Capacitors |
| SMD capacitor code      | Capacitors |
| Series/parallel (C)     | Capacitors |
| RC charge and discharge | Capacitors |
| Capacitor stored energy | Capacitors |
| Inductor color code     | Inductors  |

## DIGITAL & LOGIC CIRCUITS

|                               |                     |
|-------------------------------|---------------------|
| Karnaugh map simplification   | Logic gates         |
| I2C pull-up resistor          | Timing & interfaces |
| UART baud rate                | Timing & interfaces |
| Crystal load capacitance      | Timing & interfaces |
| Oscillator stability          | Timing & interfaces |
| PLL multiplication factor     | Timing & interfaces |
| ADC resolution / quantization | Data conversion     |
| DAC resolution                | Data conversion     |
| SNR estimation                | Data conversion     |

# Resistor color code

## What it computes

The resistance, tolerance and, on 6-band parts, temperature coefficient that a resistor's colour bands stand for — and the other way round: type a value and it shows the bands. It also says whether the value is a standard E-series value and, if not, the nearest one.

## Source

**IEC 60062:2016**, *Marking codes for resistors and capacitors*. Its table gives each colour a digit, a multiplier, a tolerance and a temperature coefficient.

## What the bands mean


A resistor is too small to print a number on, so the value is painted as rings of colour, one digit per ring. Each colour stands for a digit from 0 to 9: black 0, brown 1, red 2, orange 3, yellow 4, green 5, blue 6, violet 7, grey 8, white 9.

- **Digit bands** — the first two (4-band parts) or three (5- and 6-band parts) give the significant figures.
- **Multiplier** — the next band says how many zeros follow, as a power of ten: red is  $\times 100$ , orange  $\times 1000$ . Gold ( $\times 0.1$ ) and silver ( $\times 0.01$ ) make values below  $10\ \Omega$ .
- **Tolerance** — how far the actual resistance may be from the marked value: gold  $\pm 5\%$  means a  $1\ \text{k}\Omega$  part measures somewhere from  $950\ \Omega$  to  $1050\ \Omega$ . Leaving the band off means  $\pm 20\%$ .
- **Temperature coefficient** (6-band only) — how much the resistance drifts per degree of temperature change, in parts per million per kelvin: brown is  $100\ \text{ppm/K}$ , so  $0.01\%$  per degree.

Brown–black–red–gold is therefore 1, 0,  $\times 100$ ,  $\pm 5\%$ :  $1\ \text{k}\Omega \pm 5\%$ .

**Which end to start from.** The tolerance band usually stands apart from the others, on the right; the digits are bunched on the left. Gold and silver are never digits, so a part with gold at one end is read from the other.

**Standard values.** Resistors are made only in preferred values, the E-series, spaced so that neighbouring values just overlap within their tolerance. The tolerance band picks the one series the value is checked against — gold,  $\pm 5\%$ , means E24; brown,  $\pm 1\%$ , means E96 — and when the value is not in it, the tool gives the nearest one that is. A  $1\% 4.7\ \text{k}\Omega$  part therefore reads "Not in E96", though makers do sell it (see the E-series tool).



## How the value is worked out

4 bands:  $\text{Value} = (10 \times D1 + D2) \times \text{Multiplier}$

5, 6:  $\text{Value} = (100 \times D1 + 10 \times D2 + D3) \times \text{Multiplier}$

Typing a value runs it backwards: the value is split into as many significant digits as the band count allows, rounded, and the remaining power of ten picks the multiplier colour. The smallest value a code can show is  $0.1\ \Omega$  on 4 bands (brown–black–silver) and  $1\ \Omega$  on 5 or 6; the largest is  $99\ \text{G}\Omega$  and  $999\ \text{G}\Omega$ . Outside that there is no band for it, and the tool says so.

## Assumptions and limits

- **The 2016 tolerance colours.** IEC 60062:2016 made orange  $\pm 0.05\%$ , yellow  $\pm 0.02\%$  and grey  $\pm 0.01\%$ . Before that edition some manufacturers used grey for a non-standard  $\pm 0.05\%$ , so a grey tolerance band on an old part may mean that.
- **Colours are only as good as the light.** Red, orange and brown are easily confused on a real part, and so are violet and grey. When in doubt, measure.
- **Zero-ohm links** are marked with a single black band and are not a colour code as such.

## What it deliberately does not do

- **Capacitor and inductor colour codes** have their own tools, since the same colours mean different units and tolerances there.
- **SMD resistor codes** (the three- and four-digit and EIA-96 markings) are the SMD code tool.

## How it was checked

Against the IEC 60062:2016 table: digits, multipliers, tolerances and temperature coefficients for every colour. This check found the app had grey as  $\pm 0.05\%$ , the pre-2016 non-standard value, and no tolerance for orange or yellow; all three now follow the standard.

Brown–black–red–gold reads  $1\ \text{k}\Omega \pm 5\%$ ,  $950\ \Omega$  to  $1.05\ \text{k}\Omega$ , E6. At grey's  $\pm 0.01\%$  the range reads  $999.9\ \Omega$  to  $1.0001\ \text{k}\Omega$  — at four figures it had printed as " $1\ \text{k}\Omega$ ", which the range now shows to enough figures to see. Typing  $4.7\ \text{k}$  gives yellow–violet–red;  $0.22\ \Omega$  gives red–red–silver.

# SMD resistor code

## What it computes

The resistance printed on a surface-mount (chip) resistor, from its marking — 3-digit, 4-digit or EIA-96 — and the other way round: the marking a value would carry. It says whether the value is a standard E-series value.

## Source

**IEC 60062:2016**, *Marking codes for resistors and capacitors*, which defines the 3- and 4-character numerical codes, R as the decimal point, and the three-character code for E96 values that the industry calls EIA-96.

## What the marking means


A chip resistor is far too small for colour bands, so the value is printed as characters, three or four of them.

- **3-digit code** — usual on  $\pm 5\%$  parts. The first two digits are the value's figures, the last is how many zeros follow them. 472 is 47 followed by two zeros: 4700  $\Omega$ , 4.7 k $\Omega$ . 334 is 330 k $\Omega$ ; 100 is 10  $\Omega$ , not 100  $\Omega$ .
- **4-digit code** — usual on  $\pm 1\%$  parts, which need a third figure. The first three are the figures, the last the number of zeros. 1001 is 1.00 k $\Omega$ , 4992 is 49.9 k $\Omega$ .
- **R** — below 10  $\Omega$  (3-digit) or 100  $\Omega$  (4-digit) there is no zero count small enough, so R stands where the decimal point is. 4R7 is 4.7  $\Omega$ , R47 0.47  $\Omega$ , 47R0 47.0  $\Omega$ , R010 10 m $\Omega$  — the last one a current-sense resistor.
- **0, 000 or 0000** — a zero-ohm link, used as a jumper that a pick-and-place machine can fit.
- **EIA-96** — on small 1% parts (0603 and below) even four characters do not fit. Two digits and a letter do: the digits are not the value but its position among the 96 values of the E96 series, 01 for 100 up to 96 for 976, and the letter is the multiplier. 01C is the first E96 value, 100, times 100: 10 k $\Omega$ . 68X is 499  $\times$  0.1 = 49.9  $\Omega$ .

## How the value is worked out

3-digit:  $\text{Value} = D1D2 \times 10^{D3}$   
 4-digit:  $\text{Value} = D1D2D3 \times 10^{D4}$   
 EIA-96:  $\text{Value} = E96[\text{code}] \times \text{multiplier}$

EIA-96 multiplier letters: Z  $\times 0.001$ , Y or R  $\times 0.01$ , X or S  $\times 0.1$ , A  $\times 1$ , B or H  $\times 10$ , C  $\times 100$ , D  $\times 1000$ , E  $\times 10^4$ , F  $\times 10^5$ . Where two letters share a multiplier the tool reads both and prints the first.



Typing a value runs it backwards. A marking holds only two or three figures, so a value with more is rounded to what the part could carry — 12 345  $\Omega$  becomes 1232, 12.3 k $\Omega$  — and the tool says so. In EIA-96 only E96 values exist, so a value is moved to the nearest one: 4.7 k $\Omega$  becomes 66B, 4.75 k $\Omega$ .

## Assumptions and limits

- **The marking does not say the tolerance.** A 3-digit code is usually  $\pm 5\%$  and a 4-digit or EIA-96 code  $\pm 1\%$ , but only the datasheet or the reel label says.
- **Standard value.** The value is checked against the series of that usual tolerance: E24 for a 3-digit code, E96 for 4-digit and EIA-96. So 472 is an E24 value, and 4701 reads "Not in E96 — nearest is 4.75 k $\Omega$ ": E96 has no 4.7. A 3-digit code carries two figures, enough for every E6 to E24 value; the three figures of E48 to E192 need 4 digits or EIA-96.
- **Which scheme a part uses** shows in the marking: a letter other than R means EIA-96, and digits only are the 3- or 4-digit code, told apart by how many there are. Pick the pill that matches.
- **Small parts may not be marked at all.** 0402 and smaller usually carry no marking; the value is only on the reel.
- Some makers use their own schemes for current-sense resistors (such as a lowercase m for milliohms); those are not covered.

## What it deliberately does not do

- **Colour bands** are the Color code tool.
- **SMD capacitor and inductor markings** have their own tools, since the same characters mean picofarads or microhenries there.

## How it was checked

Against the examples given with the IEC 60062:2016 codes: 334 is 330 k $\Omega$ , 222 is 2.2 k $\Omega$ , 1001 is 1.00 k $\Omega$ , 4992 is 49.9 k $\Omega$ , R300 is 0.30  $\Omega$ , 01C is 10 k $\Omega$ ; also 68X is 49.9  $\Omega$  and 66B 4.75 k $\Omega$ .

The review found four faults, now fixed. 47 m $\Omega$  on the 3-digit pill printed "R047", four characters, where the marking is R05; 10 m $\Omega$  on the 4-digit pill printed "R0100", five characters, instead of R010, and typing R010 showed R0100 on the chip. A marking of just "0" for a zero-ohm link was rejected. And 12 345  $\Omega$  showed the code 1232 beside a resistance of 12.35 k $\Omega$ , with nothing to say the code means 12.3 k $\Omega$ .

# Resistor power rating

## What it computes

Which size a resistor needs. Enter its resistance and the voltage across it, or the current through it, and the tool gives the power it dissipates, the smallest chip size that takes it, the smallest at half its rating — the usual design margin — and the through-hole wattage class at half its rating. A chip must also be rated for the voltage across it.

Below that is the reference table: how much power each common resistor size can dissipate, with its dimensions. Surface-mount parts are listed by package code, with their maximum working voltage; through-hole parts by their wattage class, with a typical body size. It can be filtered to SMD or through-hole and searched by package or wattage.

## Source

- **SMD** — Yageo's RC\_L general-purpose thick-film chip resistor datasheet (V.11): rated power at 70 °C, maximum working voltage and operating temperature range for each size from 01005 to 2512. These are the conventional figures most tables repeat.
- **Vishay D/CRCW e3** (revision April 2026) as the counter-example: the same sizes rated higher, 0.125 W for 0603 and 0.25 W for 0805.
- **Through-hole** body sizes are typical ranges from distributor listings, not a standard.

## What it means


When current flows through a resistor, the resistor turns electrical power into heat. The power rating is the most heat it can get rid of, continuously, without being damaged or drifting out of tolerance. Exceed it and the part runs hot, its value shifts, and eventually it fails.

The power a resistor actually has to handle comes from what is across it or through it:

$$P = V^2 / R = I^2 \times R$$

A 1 kΩ resistor with 5 V across it dissipates  $25 / 1000 = 25$  mW, well inside an 0603's 100 mW. The same resistor at 12 V dissipates 144 mW: more than an 0805's 125 mW, so a 1206 at least, and a 1210 with the half-rating rule below.

**Temperature.** The rating holds only while the air around the part is at 70 °C or cooler. Above that the part can shed less heat, and the allowed power falls in a straight line to nothing at 155 °C, the maximum film temperature — 125 °C for 0201 and 01005. A common rule is to use no more than half the rating, which leaves room for a warm enclosure and for the board around the part heating it too.



**Voltage.** Each size also has a maximum working voltage, set by the gap between its terminals. For high resistance values this limit is reached before the power one:

$$V = \sqrt{P \times R}, \text{ or the maximum working voltage, whichever is lower}$$

An 0603 (1/10 W, 75 V) at 1 M $\Omega$  would reach its power rating at 316 V, but may only see 75 V, which is  $75^2 / 1 \text{ M}\Omega = 5.6 \text{ mW}$ . For high-voltage dividers the voltage rating, not the wattage, sets the size — or several resistors in series share the voltage.

## Assumptions and limits

- **Ratings depend on the series, not only the size.** The table gives the conventional general-purpose figures; anti-surge, high-power and some newer thick-film series rate the same package two or four times higher. The part's own datasheet has the last word.
- **The board matters.** An SMD rating assumes the copper pads and board carry heat away as in the maker's test. Small pads, thin copper or neighbouring hot parts lower it.
- **Through-hole sizes vary by maker and technology:** a metal-film 1/4 W part and a carbon-film one are not the same size. The ranges are for recognising a part, not for a footprint.
- Pulses are a separate question: a short pulse can exceed the continuous rating, within limits the datasheet gives as a pulse curve.
- **The size check uses the table's figures:** the Yageo ratings for chips, each with its maximum working voltage, and the wattage class for leaded parts, whose voltage rating is not in the table and depends on the part. "Half its rating" is the usual margin for a warm enclosure and hot neighbours; the full rating holds only up to 70 °C around the part, on the maker's pad layout.
- Above 200 V no chip in the table will do; several in series share the voltage, each taking its part.

## What it deliberately does not do

- **Package dimensions in detail** (pad sizes, heights) are the SMD package sizes tool.
- **Temperature rise** from a given power and board is a thermal calculation, not a rating lookup.

## How it was checked

Every SMD power rating, working voltage and the 125 °C / 155 °C limits against the Yageo RC\_L datasheet. The review corrected the 1210's size in inches, which read 0.12" × 0.125" against the 0.12" × 0.10" its name encodes, and the 0805's width to 1.25 mm; it also withdrew the claim that SMD ratings are standard across manufacturers, which the Vishay datasheet contradicts. The maximum working voltages were added then.

The size check was added in the re-walk. Its reference cases: 1 k $\Omega$  at 12 V is 144 mW, a 1206 at full rating, a 1210 or a 1/2 W leaded part at half; a 0.1  $\Omega$  shunt at 2 A is 0.4 W, a 1210 or at half a 2512; 1 M $\Omega$  at 100 V is only 10 mW but needs an 0805, the first chip rated over 75 V. The re-walk also corrected this section's own example, which had said 144 mW fits "an 0805 at least" when an 0805 is rated 125 mW.

# E-series values

## What it computes

The standard value nearest to the one you want, in the series you choose (E6 to E192), how far off it is in percent, and the whole series across the decade of your value, with the nearest one highlighted.

## Source

**IEC 60063:2015**, *Preferred number series for resistors and capacitors*. It lists every series value by value, and its tolerances: E6  $\pm 20\%$ , E12  $\pm 10\%$ , E24  $\pm 5\%$ , E48  $\pm 2\%$ , E96  $\pm 1\%$ , E192  $\pm 0.5\%$  and better.

## What it means

Resistors and capacitors are not made in every value. It would be impossible to stock them all, and pointless: a  $\pm 5\%$  1000  $\Omega$  resistor can be anything from 950 to 1050  $\Omega$ , so a separate 1010  $\Omega$  part would overlap it completely. Instead, manufacturers make a fixed set of **preferred values**, the E-series.

Each series splits a decade — 1 to 10, 10 to 100, and so on — into a fixed number of steps: 6 for E6, 24 for E24, 96 for E96. The steps are not equal in ohms but equal in **ratio**, so each value is the same percentage above the previous one, whatever the decade. And the number of steps is chosen so that the gap between neighbours matches the tolerance: in E24 each value is about 10% above the last, so  $\pm 5\%$  parts either side just meet. That is why the tolerance tells you the series —  $\pm 5\%$  parts come in E24 values,  $\pm 1\%$  in E96.


In practice: pick the tolerance you need, look up the nearest value in its series, and check whether the percent difference matters for your circuit. If it does, a tighter series has a closer value, or two parts in series or parallel can make up the difference.

## How the values are worked out

$$E_i = 10^{(i / N)}, \quad i = 0 \dots N-1, \quad \text{rounded to 2 figures (E6-E24) or 3 (E48-E192)}$$

N is the number of values in the series. E48, E96 and E192 follow this formula, with one exception: the standard lists 920 in E192 where the formula gives 919. E6, E12 and E24 are older than the formula and were chosen by hand; eight E24 values differ from it — 2.7, 3.0, 3.3, 3.6, 3.9, 4.3, 4.7 and 8.2 — which is why 4.7 k $\Omega$ , not 4.6 k $\Omega$ , is the value everyone knows. The tool uses the standard's lists, not the formula, for those.

The nearest value is the one with the smallest percentage difference from what you typed. The top of a decade counts as the start of the next one, so 97 k $\Omega$  in E24 gives 100 k $\Omega$ , not 91 k $\Omega$ .



## Assumptions and limits

- **E3** ( $\pm 40\%$ : 1.0, 2.2, 4.7) exists in the standard but is rarely used for new parts and is not offered.
- **Exactly between two values** the lower one is given; both are equally far off.
- **One tolerance, one series.** Throughout the app a tolerance picks exactly one series:  $\pm 1\%$  means E96, and the sliders and hints step through E96 only. Manufacturers also sell the E24 values at  $\pm 1\%$  — Yageo's RC\_L datasheet lists "1% (E24/E96)" — so a 1% 4.7 k $\Omega$  part is easy to buy even though the app suggests 4.75 k $\Omega$ . The simpler rule is kept on purpose.
- **Capacitors** follow the same series, but in practice ceramic capacitors are mostly stocked in E6 or E12 values, and electrolytics in E3 or E6.
- The table shows the decade of the nearest value, so the value found is always highlighted in it: for 97 k $\Omega$  in E24 that is the 100 k $\Omega$  decade.

## What it deliberately does not do

- **Combining two standard values** to hit an exact one is a series/parallel question, the Series/parallel tool's job.

## How it was checked

Every E48 and E96 value against the IEC 60063 list, and E192 against the same list. The review found the E192 exception missing: the tool showed 919 where the standard has 920, so 9.2 k $\Omega$  was reported as not standard; it now uses 920. The E6, E12 and E24 lists, including the eight values that differ from the formula, match the standard. The table also showed the decade of the typed value, so a nearest value in the next decade (97 k $\Omega$  → 100 k $\Omega$ ) was not highlighted; it now shows the nearest value's decade.

# Series/parallel (R)

## What it computes

The total resistance of two to four resistors connected in series or in parallel, the range it can take given the parts' tolerance, and the nearest single standard part to that total. Each resistor has a slider that steps through standard values.

## Source

Ohm's law and Kirchhoff's laws; no standard is needed beyond them. The standard values come from IEC 60063, as in the E-series tool.

## What it means

**In series** — resistors connected end to end, so the same current has to pass through each of them in turn. Each one drops a voltage, and the drops add up, so the resistances add:

$$R = R1 + R2 + R3 + \dots$$

The total is always more than the largest one.

**In parallel** — resistors connected side by side between the same two points, so each sees the same voltage and the current divides between them. Each gives the current another path, so the total is always less than the smallest one. It is the conductances,  $1/R$ , that add:


$$\begin{aligned} 1/R &= 1/R1 + 1/R2 + 1/R3 + \dots \\ R &= R1 \times R2 / (R1 + R2) \quad \text{for two} \end{aligned}$$

Two equal resistors in parallel give half of one; ten in parallel, a tenth. A resistor much larger than another barely changes it: 1 kΩ in parallel with 1 MΩ is 999 Ω.

**Why combine them.** The common reason is to get a value the E-series does not stock from two it does — 1 kΩ in parallel with 4.7 kΩ is 824.6 Ω, within a tenth of a percent of the E96 value 825 Ω. Others: to share the power or the voltage between several parts, each within its own rating.

## Tolerance of the total

If every resistor is off by the same fraction in the same direction, the series sum and the parallel combination are both off by exactly that fraction. So in the worst case the total has the same tolerance as the parts, no better and no worse, and that is the range the tool shows. In practice the errors of separate parts are



independent and partly cancel, so the total is usually closer than that — but a worst-case design should not count on it.

## Assumptions and limits

- **Up to four resistors**, all in series or all in parallel. Mixed networks (two in series, in parallel with a third) are worked out in steps: combine one group, then use its total as a single resistor.
- **A zero-ohm resistor in parallel** shorts the group, and the total is 0.
- Wire and contact resistance are ignored, which matters only for milliohm values.

## What it deliberately does not do

- **Capacitors and inductors** combine the other way round and have their own tools.
- **Finding the best pair** for a target value automatically is not done; the sliders make trying pairs quick.

## How it was checked

By hand:  $1\text{ k} + 2.2\text{ k} = 3.2\text{ k}\Omega$ ;  $1\text{ k} \parallel 1\text{ k} = 500\text{ }\Omega$ ;  $1\text{ k} \parallel 4.7\text{ k} = 824.6\text{ }\Omega$ , nearest E96  $825\text{ }\Omega$ ; a third  $1\text{ k}\Omega$  in parallel gives  $451.9\text{ }\Omega$ ;  $1\text{ k} \parallel 0 = 0\text{ }\Omega$ . The tolerance ranges are the total  $\times (1 \pm t)$ .

# Voltage divider

## What it computes

For two resistors in series across a voltage, any one of the output voltage, the top resistor R1 or the bottom resistor R2, from the other three. It also gives the output's worst-case spread from the resistors' tolerance, the current through the pair and the power in each resistor, and when a resistor is solved for, the nearest standard value and the output that value would give.

## Source

Ohm's law applied to two resistors in series; standard values from IEC 60063.

## What it means

A voltage divider makes a smaller voltage out of a larger one. Two resistors are connected in series from the input voltage to ground. The same current flows through both, so each drops a share of the voltage in proportion to its resistance, and the point between them — the output — sits at a fixed fraction of the input:

$$V_{out} = V_{in} \times R2 / (R1 + R2)$$

Equal resistors give half. A 10 kΩ over a 4.7 kΩ gives 4.7 / 14.7, about a third. Typical uses: bringing a 12 V signal down to what a 3.3 V ADC can read, setting a reference voltage, or biasing a transistor.

Solved for either resistor:


$$R1 = R2 \times (V_{in} - V_{out}) / V_{out}$$

$$R2 = R1 \times V_{out} / (V_{in} - V_{out})$$

**Choosing the values.** Only the ratio sets Vout, so 1 kΩ/1 kΩ and 1 MΩ/1 MΩ both halve the voltage. What the size decides is the current: small values waste power (12 V across 2 kΩ is 6 mA, 72 mW), large ones make the output weak, easily pulled down by whatever it feeds and more sensitive to noise. Tens of kΩ is a common middle ground for signals.

## Tolerance

The output is furthest off when the two resistors miss in opposite directions: R2 high and R1 low pushes Vout up, the reverse pulls it down.



$$V_{out \max} = V_{in} \times R_2(1+t) / (R_1(1-t) + R_2(1+t))$$

$$V_{out \min} = V_{in} \times R_2(1-t) / (R_1(1+t) + R_2(1-t))$$

So a divider built from  $\pm 1\%$  parts is not a  $\pm 1\%$  divider: 10 k $\Omega$  over 4.7 k $\Omega$  at  $\pm 1\%$  spreads by  $\pm 1.37\%$ . The spread is widest when the output is a small fraction of the input, and approaches  $\pm 2t$ .

## Loading

The formula assumes nothing draws current from the output. Anything that does — an ADC input, a transistor base, a meter — is a resistor in parallel with  $R_2$ , and lowers the output. Seen from the output, the divider behaves like a source with an internal resistance of  $R_1 \parallel R_2$ , so a load  $R_L$  lowers  $V_{out}$  by roughly

$$(R_1 \parallel R_2) / R_L$$

1% for a load a hundred times  $R_1 \parallel R_2$ , 10% for one ten times. The fix is lower divider resistances or a buffer (the op-amp voltage follower).

## Assumptions and limits

- **Unloaded**, as above.
- **Positive voltages**. For a negative supply, enter both voltages as positive values; the ratio and the resistors are the same.
- $V_{out}$  must be between 0 and  $V_{in}$  when solving for a resistor; a divider can only divide.
- The standard value offered is from the one series the chosen tolerance implies ( $\pm 1\% \rightarrow E96$ ).

## What it deliberately does not do

- **A load resistor** is not an input; the loading rule above covers it.
- **Adjustable dividers** (a potentiometer with fixed end resistors) are not modelled.

## How it was checked

By hand: 12 V over 10 k/10 k gives 6 V, 5.94–6.06 V at  $\pm 1\%$ , 600  $\mu$ A, 3.6 mW each. 10 k over 4.7 k gives 3.837 V, 3.78–3.89 V ( $\pm 1.37\%$ ).  $R_1$  for 12 V to 3.3 V with  $R_2 = 10$  k $\Omega$  is 26.36 k $\Omega$ , nearest E96 26.1 k $\Omega$  giving 3.32 V.  $R_2$  for 5 V to 3.3 V with  $R_1 = 10$  k $\Omega$  is 19.41 k $\Omega$ , E96 19.6 k $\Omega$  giving 3.31 V.  $V_{out}$  equal to  $V_{in}$ , or below zero, is refused with a reason.

# Current divider

## What it computes

For a current flowing into two resistors in parallel, any one of the current in the first branch  $I_1$ ,  $R_1$  or  $R_2$ , from the other three. It also gives the current in the second branch, the voltage across both, the spread of  $I_1$  from the resistors' tolerance, and when a resistor is solved for, the nearest standard value and the current it would give.

## Source

Ohm's law and Kirchhoff's current law; standard values from IEC 60063.

## What it means

When a current reaches two resistors connected in parallel, it splits: part goes through each branch and the two parts add back up to the whole (Kirchhoff's current law). Both branches have the same voltage across them, so by Ohm's law each takes a current inversely proportional to its resistance — the easier path takes more:

$$I_1 = I_{in} \times R_2 / (R_1 + R_2)$$

$$I_2 = I_{in} \times R_1 / (R_1 + R_2)$$

This is the voltage divider turned inside out, and the place people trip: the current in branch 1 is set by the *other* resistor,  $R_2$ . With 20 mA into 1 kΩ in parallel with 2 kΩ, the 1 kΩ branch takes two thirds, 13.3 mA, and the 2 kΩ branch one third, 6.7 mA.

Solved for a resistor:


$$R_1 = R_2 \times (I_{in} - I_1) / I_1$$

$$R_2 = R_1 \times I_1 / (I_{in} - I_1)$$

**Where it is used.** Shunts that extend a meter's range (a small resistor in parallel takes most of the current, the meter a known fraction), sharing current between parallel parts, and understanding why a low-value resistor in parallel with a circuit steals most of its current.

## Tolerance

As with the voltage divider,  $I_1$  is furthest off when the two resistors miss in opposite directions, so the spread of  $I_1$  can exceed the parts' own tolerance. When  $I_1$  is the larger share, it is less sensitive: 1 k || 2 k at ±1% moves  $I_1$  by only ±0.66%, because most of the error lands in the smaller branch current.



## Assumptions and limits

- **The total current is taken as given**, as if from a current source. In a real circuit fed from a voltage,  $I_{in}$  itself depends on  $R_1 \parallel R_2$ .
- **Two branches**. For more, combine all but one into a single resistance first with the Series/parallel tool.
- $I_1$  must be between 0 and  $I_{in}$  when solving for a resistor.
- The standard value offered is from the one series the chosen tolerance implies.

## What it deliberately does not do

- **Power in each branch** is not shown; the voltage across the branches is, and  $P = V^2 / R$  from there.

## How it was checked

By hand: 20 mA into  $1\text{ k} \parallel 1\text{ k}$  is 10 mA each at 10 V; into  $1\text{ k} \parallel 2\text{ k}$ , 13.33 mA and 6.67 mA at 13.3 V,  $I_1$  between 13.2 and 13.4 mA at  $\pm 1\%$ .  $R_1$  for 15 mA of 20 mA with  $R_2 = 1\text{ k}\Omega$  is  $333.3\ \Omega$ , E96 332  $\Omega$ .  $R_2$  for 5 mA of 20 mA with  $R_1 = 1\text{ k}\Omega$  is  $333.3\ \Omega$ .  $I_1$  above  $I_{in}$  is refused. The review labelled the voltage in the result line, which had been a bare "10 V" with nothing to say what it was.

# Wheatstone bridge

## What it computes

For a balanced Wheatstone bridge, any one of its four resistors from the other three. Solving for the unknown  $R_x$ , it also gives the ratio  $R_2 / R_1$  that multiplies  $R_3$  into  $R_x$ ; solving for one of the fixed arms, the nearest standard value.

## Source

The balance condition follows from Ohm's law applied to the two dividers that make up the bridge. The circuit is Samuel Hunter Christie's (1833), popularised by Charles Wheatstone.

## What it means

A Wheatstone bridge measures an unknown resistance by comparing it with known ones, instead of measuring a current and a voltage.

Four resistors form a diamond. A supply  $V$  is connected across the top and bottom corners, so the left side ( $R_1$  over  $R_3$ ) and the right side ( $R_2$  over  $R_x$ ) are two voltage dividers on the same supply. A sensitive meter — a galvanometer,  $G$  — joins the midpoints of the two sides.

If the two midpoints are at the same voltage, no current flows through the meter and it reads zero. The bridge is then **balanced**, and that happens when both dividers have the same ratio:

$$R_1 / R_2 = R_3 / R_x \quad \text{so} \quad R_x = R_2 \times R_3 / R_1$$


The supply voltage and the meter's own characteristics drop out entirely: the only thing that matters is whether the meter reads zero, which even a crude meter can tell very precisely. That is what made the bridge the standard way to measure resistance accurately.

**In use.**  $R_1$  and  $R_2$  are fixed and set the range: with  $R_2 / R_1 = 10$ ,  $R_x$  is ten times  $R_3$ .  $R_3$  is a calibrated adjustable resistor. Turn  $R_3$  until the meter reads zero, and  $R_x = 10 \times R_3$ .

The same circuit, left unbalanced, is how strain gauges, pressure sensors and many thermometers are read: a small change in one arm produces a small voltage across the meter's position, measured by an amplifier.

## Assumptions and limits

- **Balance only.** The tool gives the resistance that balances the bridge; it does not compute the voltage or current through the meter when the bridge is off balance.



- **The precision is that of the known arms.** Rx can be no more accurate than R1, R2 and R3 combined; with  $\pm 0.1\%$  parts, roughly  $\pm 0.3\%$  worst case.
- The arm used as a divisor cannot be zero: R1 when solving Rx, Rx when solving R1, R3 for R2 and R2 for R3.

## What it deliberately does not do

- **Unbalanced output** for sensor bridges (the voltage across G as a function of a changing arm) is a different calculation, not done here.
- AC bridges for capacitance and inductance (Maxwell, Schering, Wien) are not covered.

## How it was checked

By hand: R1 = 1 k, R2 = 2 k, R3 = 3 k balance with Rx = 6 k $\Omega$ , ratio 2; each of the other three arms solved from the remaining ones gives back 1 k, 2 k and 3 k. R1 = 0 is refused. The review replaced the "nearest standard value" once shown for Rx, which is the resistance being measured rather than a part to buy, with the bridge ratio; dropped "(unknown)" from Rx's label, which stayed even when Rx was an input; and made the supply's V label neutral, which had taken R1's colour.

# Delta-Y transform

## What it computes

Converts three resistors connected as a triangle (Delta,  $\Delta$ ) into the three of an equivalent star (Y), or the reverse, so that the two behave identically between their three terminals.

## Source

Arthur Edwin Kennelly's transform (1899), derived from requiring the resistance between each pair of terminals to be the same in both networks.

## What it means

Three resistors can be connected to three terminals A, B and C in two ways:

- **Delta ( $\Delta$ )** — a triangle. One resistor between each pair of terminals:  $R_{ab}$ ,  $R_{bc}$ ,  $R_{ca}$ .
- **Y (star, also called T)** — three resistors meeting at a common centre point, one leg to each terminal:  $R_a$ ,  $R_b$ ,  $R_c$ .

Seen only from the three terminals, a Delta and a suitably chosen Y cannot be told apart: the resistance between any two terminals is the same, whatever else is connected. So one can always be replaced by the other.


**Why that is useful.** Many networks are neither series nor parallel anywhere — the classic case is a bridge, five resistors in a diamond with one across the middle. Replacing one of its triangles by the equivalent star turns it into ordinary series and parallel combinations that can be simplified step by step. The same transform turns up in three-phase power, where loads are wired in delta or in star.

## The formulas

**Delta to Y** — each leg of the star is the product of the two Delta sides that touch its terminal, divided by the sum of all three:

$$\begin{aligned} R_a &= R_{ab} \times R_{ca} / (R_{ab} + R_{bc} + R_{ca}) \\ R_b &= R_{ab} \times R_{bc} / (R_{ab} + R_{bc} + R_{ca}) \\ R_c &= R_{bc} \times R_{ca} / (R_{ab} + R_{bc} + R_{ca}) \end{aligned}$$

**Y to Delta** — each side of the triangle is the sum of the products of the legs taken in pairs, divided by the leg *opposite* that side:



$$R_{ab} = (R_a \cdot R_b + R_b \cdot R_c + R_c \cdot R_a) / R_c$$

$$R_{bc} = (R_a \cdot R_b + R_b \cdot R_c + R_c \cdot R_a) / R_a$$

$$R_{ca} = (R_a \cdot R_b + R_b \cdot R_c + R_c \cdot R_a) / R_b$$

A way to remember them: Delta to Y divides, so the star's legs are smaller; Y to Delta multiplies, so the triangle's sides are larger. With three equal parts the Y is a third of the Delta: a 300  $\Omega$  triangle is a 100  $\Omega$  star.

## Assumptions and limits

- **All three must be greater than zero.** A zero-ohm side or leg joins two terminals directly, and the network is no longer a true three-terminal one.
- **Resistance only.** The transform holds just as well for impedances (capacitors and inductors, in AC circuits), but this tool works in ohms.

## What it deliberately does not do

- **Solving a whole network** (such as the full bridge) is not automated; the tool does the one step that makes the rest ordinary.

## How it was checked

By hand: a 300  $\Omega$  Delta gives a 100  $\Omega$  Y, and back.  $R_{ab} = 10$ ,  $R_{bc} = 20$ ,  $R_{ca} = 30$   $\Omega$  give  $R_a = 5$ ,  $R_b = 3.333$ ,  $R_c = 10$   $\Omega$ , and those three converted back give 10, 20 and 30  $\Omega$ .

# NTC/PTC thermistor

## What it computes

For a temperature-sensing resistor, its resistance at a given temperature or the temperature for a measured resistance, from the datasheet's reference resistance and coefficient; and its sensitivity, in percent per degree, at that point. A curve from  $-20$  to  $100$  °C shows where the point sits.

- **NTC** — the B-parameter model, used for ordinary NTC thermistors.
- **PTC** — a linear model, for silicon (KTY) sensors and, approximately, platinum RTDs.

## Source

- The B-parameter equation, as NTC datasheets define it. The Vishay NTCLE100E3 datasheet defines B as  $\ln(R_1/R_2) / (1/T_1 - 1/T_2)$  and gives  $B_{25/85} = 3977$  K for its 10 kΩ part, with a full resistance-temperature table, which is what the tool was checked against.
- The linear PTC model is the definition of a temperature coefficient  $\alpha$ .


## What it means

A **thermistor** is a resistor whose resistance changes strongly with temperature, made to be used as a temperature sensor. Put it in a voltage divider, read the voltage with an ADC, work out the resistance, and from that the temperature.

- **NTC** (negative temperature coefficient): resistance falls as temperature rises, steeply — about 4% per °C near room temperature for a typical part. A 10 kΩ NTC is 10 kΩ at 25 °C, about 33 kΩ at 0 °C and about 1 kΩ at 85 °C.
- **PTC** (positive temperature coefficient): resistance rises with temperature. The sensors modelled here rise steadily; a KTY silicon sensor by about 0.7% per °C, a Pt100 platinum RTD by 0.385% per °C.

The datasheet gives two numbers that define the part:

- **R0 at T0** — its resistance at a reference temperature, almost always 25 °C. "10 kΩ NTC" means 10 kΩ at 25 °C.
- **B** (for an NTC), in kelvin — how steep the curve is. A higher B means a larger change per degree. Or  $\alpha$  (for a linear PTC) in percent per degree.



## The formulas

$$\begin{aligned}\text{NTC: } R &= R_0 \times e^{(B \times (1/T - 1/T_0))} \\ T &= 1 / (1/T_0 + \ln(R/R_0) / B) \\ \text{PTC: } R &= R_0 \times (1 + \alpha \times (T - T_0))\end{aligned}$$

T and T<sub>0</sub> are in kelvin (°C + 273.15) in the NTC equations. The sensitivity is the slope of the curve as a percentage:  $-B / T^2$  for the NTC, which grows steeper the colder it is, and simply  $\alpha$  for the linear PTC.

## How good the B model is

B is measured between two temperatures, written as a suffix: **B25/85** is worked out from the resistance at 25 °C and at 85 °C, B25/50 from 25 and 50 °C. The model then goes exactly through those two points and bends away from the real curve outside them. Checked against the Vishay 10 kΩ part's own table, with B25/85 = 3977 K:

| Temperature | Datasheet | B model | Difference         |
|-------------|-----------|---------|--------------------|
| −40 °C      | 332.1 kΩ  | 412 kΩ  | +24%, about 3.5 °C |
| −20 °C      | 96.4 kΩ   | 107 kΩ  | +11%, about 2 °C   |
| 0 °C        | 32.6 kΩ   | 33.9 kΩ | +4%, under 1 °C    |
| 85 °C       | 1.07 kΩ   | 1.07 kΩ | exact              |
| 100 °C      | 677 Ω     | 685 Ω   | +1%                |

So: use the B value whose temperature pair covers the range you measure, and for accuracy well below 0 °C use the datasheet's table or the Steinhart-Hart equation, which fits three points instead of two.

## Assumptions and limits

- **Self-heating is ignored.** The current used to measure a thermistor also warms it. Datasheets give a dissipation factor (a few mW/°C); keep the measuring current small enough that it adds a fraction of a degree.
- **The part's own tolerance** — typically ±1 to ±5% on R<sub>25</sub> and ±0.5 to ±3% on B — adds to the model error.
- **Switching PTCs** (resettable polymer fuses, ceramic PTC heaters and motor protectors) jump sharply by orders of magnitude at a trip temperature. No single smooth formula describes them, and they are not modelled.
- A platinum RTD is not exactly linear; the linear model is within a few tenths of a degree over 0–100 °C, and the full Callendar-Van Dusen equation is needed beyond that.

## What it deliberately does not do

- **Steinhart-Hart** three-coefficient fitting is not offered.

- **The divider around the sensor** (choosing the fixed resistor for best resolution) is the Voltage divider tool's arithmetic.

## How it was checked

Against the Vishay NTCLE100E3 datasheet table as above. At 85 °C the model gives 1.07 k $\Omega$ , as the datasheet does, and typing 1.07 k $\Omega$  back gives 85.0 °C; the sensitivity at 25 °C is  $-4.47\text{ \%/}^{\circ}\text{C}$ . The linear PTC with  $R_0 = 1\text{ k}\Omega$ ,  $\alpha = 0.7\text{ \%/}^{\circ}\text{C}$  gives 1.525 k $\Omega$  at 100 °C.

The review made the temperature field accept negative values on an iPhone, whose number keypad has no minus key; limited the resistance units to  $\Omega$ , k $\Omega$  and M $\Omega$ ; and replaced the table of "typical" B values, which had no source, with the Vishay figures.

# Ceramic cap code

## What it computes

The capacitance and tolerance printed on a ceramic capacitor as a three- or four-digit code with an optional letter, and the other way round: the marking for a value. It says whether the value is a standard E-series value.


## Source

- **EIA-198** (RS-198), the code for ceramic capacitors: digits in picofarads, with 8 and 9 as the multipliers  $\times 0.01$  and  $\times 0.1$  for values under 10 pF.
- **IEC 60062:2016**, for R as the decimal point and the tolerance letters.

## What the marking means

Ceramic disc and small film capacitors have no room for "100 nF", so the value is printed as a code, always counted in **picofarads** (pF). 1 nF is 1000 pF, and 1  $\mu$ F is 1 000 000 pF.

- **Three digits** — the first two are the value's figures, the last is how many zeros follow them. 104 is 10 followed by four zeros: 100 000 pF, which is 100 nF or 0.1  $\mu$ F — the most common capacitor there is. 472 is 4700 pF, 4.7 nF. 220 is 22 pF, not 220 pF.
- **Small values** — below 10 pF there is no zero count small enough, so the last digit 9 means  $\times 0.1$  and 8 means  $\times 0.01$ : 479 is  $47 \times 0.1 = 4.7$  pF, 109 is 1.0 pF. Some makers write R for the decimal point instead: 4R7 is 4.7 pF.
- **Four digits** — three figures and a multiplier, for closer values: 1002 is 10 000 pF, 10 nF.
- **A letter after the digits** is the tolerance:



| Letter | 10 pF and below | Above 10 pF  |
|--------|-----------------|--------------|
| B      | $\pm 0.1$ pF    | $\pm 0.1\%$  |
| C      | $\pm 0.25$ pF   | $\pm 0.25\%$ |
| D      | $\pm 0.5$ pF    | $\pm 0.5\%$  |
| F      | $\pm 1$ pF      | $\pm 1\%$    |
| G      | $\pm 2$ pF      | $\pm 2\%$    |
| J      |                 | $\pm 5\%$    |
| K      |                 | $\pm 10\%$   |
| M      |                 | $\pm 20\%$   |
| Z      |                 | +80% / -20%  |

Small capacitors use a tolerance in picofarads because a percentage of a few picofarads would be a fraction of the stray capacitance of the leads. Z is typical of cheap high-capacitance ceramics, whose value is mostly guaranteed not to be too low.

So **473K** is 47 nF  $\pm 10\%$ , **104M** 100 nF  $\pm 20\%$ , **220J** 22 pF  $\pm 5\%$ .

**Other markings on the part.** A code such as **X7R**, **X5R**, **C0G** or **NP0** is the dielectric, which says how the capacitance changes with temperature and voltage — C0G/NP0 barely at all, X7R by  $\pm 15\%$  over  $-55$  to  $+125$  °C, Y5V by much more. A number followed by V, or a code such as 1H or 2A, is the voltage rating. Neither is part of the value code.

## Assumptions and limits

- **The value is nominal.** Class 2 ceramics (X7R, X5R, Y5V) lose capacitance with DC voltage applied: a 10  $\mu$ F X5R part can be well under half of that at its rated voltage. The code says nothing about it; the datasheet's DC-bias curve does.
- **SMD ceramic chips are almost never marked:** they are too small, and the value is only on the reel. The code is mostly seen on through-hole discs and on film capacitors.
- A value typed in is encoded with R below 10 pF (4R7); the 479 form means the same and is read as well.
- **Standard value.** The tolerance letter picks the one series the value is checked against: J ( $\pm 5\%$ ) E24, K ( $\pm 10\%$ ) E12, M ( $\pm 20\%$ ) E6, as for resistors. With no letter, or Z, it is E6, the series most ceramics are stocked in. At 10 pF and below, B to G give a tolerance in picofarads, which implies no series, and the tool says so.

## What it deliberately does not do

- **Film capacitor markings** such as 4n7 and the voltage codes have the Film code tool; **SMD capacitor** two-character codes have their own tool.

## How it was checked

104 → 100 nF, 473K → 47 nF  $\pm 10\%$ , 220J → 22 pF  $\pm 5\%$ , 106M → 10  $\mu$ F  $\pm 20\%$ , 4R7 → 4.7 pF, 8R2C → 8.2 pF  $\pm 0.25$  pF.

The review found three faults, now fixed. Codes ending in 8 or 9 were read as powers of ten, so 479 came out as 47 mF instead of 4.7 pF; they now follow EIA-198. B, C and D above 10 pF were read as picofarads, so 101D showed  $\pm 0.5$  pF instead of  $\pm 0.5\%$ . And the units offered ran to millifarads and farads, which no ceramic capacitor reaches; they are now pF, nF and  $\mu$ F. The drawing now shows the marking as typed (479 stays 479 rather than becoming 4R7).

The re-walk checked the value against the tolerance's own series rather than the coarsest one it happened to be in: 182 is "Not in E6 — nearest is 1.5 nF", 182K an E12 value, 479C "no E-series implied".

# Film capacitor code

## What it computes

The capacitance and tolerance printed on a film capacitor, in either of the two ways film parts are marked — the 3-digit picofarad code or the letter code (4n7) — and the other way round: the marking for a value.

## Source

- **EIA-198** for the 3-digit code, the same as on ceramic capacitors.
- **IEC 60062:2016** for the letter code (the "RKM" code: p, n,  $\mu$  in place of the decimal point) and the tolerance letters.

## What the marking means

Film capacitors — polyester (PET), polypropylene (PP) and similar, usually a small rectangular box with two leads — carry their value in one of two ways.

**The 3-digit code**, in picofarads. The first two digits are the figures, the last how many zeros follow: 104 is 100 000 pF = 100 nF, 473 is 47 nF, 102 is 1 nF. As on ceramics, a last digit of 9 or 8 means  $\times 0.1$  or  $\times 0.01$  for values under 10 pF, though film parts that small are rare.


**The letter code.** The unit letter stands where the decimal point would be:

- p = picofarad, n = nanofarad,  $\mu$  (often written u) = microfarad
- 4n7 = 4.7 nF, n33 = 0.33 nF (330 pF),  $\mu 1$  or u1 = 0.1  $\mu$ F (100 nF), 2 $\mu$ 2 = 2.2  $\mu$ F, 100p = 100 pF

It is the same idea as 4k7 for a 4.7 k $\Omega$  resistor: no decimal point to rub off or be misread.

**The tolerance letter** that often follows either form:

| Letter | Tolerance   |
|--------|-------------|
| F      | $\pm 1\%$   |
| G      | $\pm 2\%$   |
| H      | $\pm 2.5\%$ |
| J      | $\pm 5\%$   |
| K      | $\pm 10\%$  |
| M      | $\pm 20\%$  |



**The voltage code.** Many film capacitors add a two-character voltage code in front of the value, per the same EIA scheme: 1H = 50 V, 2A = 100 V, 2E = 250 V, 2G = 400 V, 2J = 630 V. So **2A104J** is a 100 V, 100 nF,  $\pm 5\%$  part. Others print the voltage plainly (63 V, 400 V=).

## Assumptions and limits

- **Plain decimals** such as ".1" or "0.1" on older and larger parts mean microfarads (0.1  $\mu\text{F}$ ); the tool reads only the two coded forms.
- **The voltage code is not read** from the marking; type the part of the marking that follows it.
- A value typed in is encoded with the unit that gives the shortest marking; n33 and 330p mean the same.
- **Standard value.** The tolerance letter picks the one series the value is checked against: J ( $\pm 5\%$ ) E24, K ( $\pm 10\%$ ) E12, M ( $\pm 20\%$ ) E6, G and H E48, F E96. With no letter it is E12, since film parts are mostly  $\pm 10\%$  and stocked in E12 values.

## What it deliberately does not do

- **Ceramic codes with the pF-based tolerance letters** (B, C, D) are the Ceramic code tool; film capacitors are not made small enough to need them.
- **Electrolytic** capacitors print their value and voltage in plain text.

## How it was checked

104J  $\rightarrow$  100 nF  $\pm 5\%$ , 473K  $\rightarrow$  47 nF  $\pm 10\%$ , 102H  $\rightarrow$  1 nF  $\pm 2.5\%$ , 4n7  $\rightarrow$  4.7 nF, 4N7J  $\rightarrow$  4.7 nF  $\pm 5\%$ , n33  $\rightarrow$  330 pF, u1 and  $\mu 1 \rightarrow$  100 nF, 2u2K  $\rightarrow$  2.2  $\mu\text{F}$   $\pm 10\%$ , 100p  $\rightarrow$  100 pF.

The review found the tolerance table borrowed from the inductor tool, which had no H ( $\pm 2.5\%$ ), a common film tolerance; added it. It made the unit letter case-blind (4N7 was rejected), applied the 8 and 9 multipliers and the R-code length fix from the ceramic tool, limited the units to pF, nF and  $\mu\text{F}$ , and made the drawing show the marking as typed.

The re-walk checked the value against the tolerance's own series: 562 is an E12 value, 912J an E24 one, 562M "Not in E6 — nearest is 4.7 nF". A value read from a marking now shows in its own unit, n33 as 330 pF.

# SMD capacitor code

## What it computes

The capacitance printed on a surface-mount capacitor, usually a tantalum part, with the letter printed next to it read either as the rated voltage or as the tolerance; and the other way round, the marking for a value.

## Source

- The capacitance digits are the EIA-198 picofarad code, as on ceramic parts.
- The marking layout and voltage letters from the **Kyocera-AVX TAJ** tantalum datasheet: "227 A" is 220  $\mu$ F at 10 V (A = 10 V) on its larger cases; on its small cases the letter sits above the digits, "J" over "106" for 10  $\mu$ F at 6.3 V. The rest of the voltage letters are the EIA tantalum code.
- The **KEMET T491** datasheet lists what its marking holds: polarity band, picofarad code, rated voltage and a date code.

## What the marking means

Most small ceramic chip capacitors carry no marking at all — they are too small to print on, and the value is only on the reel. The SMD capacitors that are marked are mostly **tantalum** (and polymer) capacitors, which are larger and come in moulded cases.

**The digits** are the capacitance in picofarads, as on a ceramic disc: the last digit is how many zeros follow the first two. 106 is 10 000 000 pF = 10  $\mu$ F; 227 is 220  $\mu$ F; 475 is 4.7  $\mu$ F.

**The letter** is, on most tantalum parts, the **rated voltage** — the most it may see continuously:

**Capacitor SMD code**

3 digit code + voltage or tolerance letter

106

Voltage letter    Tolerance letter

**MARKING ON THE CHIP**

Code  
106

**OR ENTER A VALUE**

Capacitance  
10  $\mu$ F

Voltage letter  
None

**RESULTS**

Capacitance Calculated  
**10  $\mu$ F**  
E6 standard value  
No voltage letter

**FORMULA**

Value = (D1D2)  $\times 10^{D3}$ , in pF

Printed SMD capacitors are mostly tantalum: small ceramic chips are too small to mark. The digits are the usual code in picofarads, the last one being how many zeros follow: 106 is 10  $\mu$ F, 227 is 220  $\mu$ F. A letter next to them, before or after, is usually the rated voltage: G 4 V, J 6.3 V, A 10 V, C 16 V, D 20 V, E 25 V, V 35 V, T 50 V, so 227A is 220  $\mu$ F at 10 V. Some parts use a tolerance letter

Home    Search    Tools    Favorites

| Letter | Rated voltage |
|--------|---------------|
| G      | 4 V           |
| J      | 6.3 V         |
| A      | 10 V          |
| C      | 16 V          |
| D      | 20 V          |
| E      | 25 V          |
| V      | 35 V          |
| T      | 50 V          |

It may be printed after the digits (227A) or before them (J106, often on a line of its own above); the tool reads both. Some parts instead use a **tolerance** letter (K  $\pm 10\%$ , M  $\pm 20\%$ ). The two tables share letters — G and J mean 4 V and 6.3 V in one,  $\pm 2\%$  and  $\pm 5\%$  in the other — so the marking alone cannot say which is meant; the datasheet can, and the pill chooses.

**Polarity.** A tantalum capacitor is polarised: the band or bevel at one end marks the **positive** terminal — the opposite of an aluminium electrolytic, whose stripe marks the negative. Fitted backwards, a tantalum can fail short and burn.

## Assumptions and limits

- **Derate the voltage.** Tantalum ( $\text{MnO}_2$ ) capacitors are usually run at no more than half their rated voltage, to survive turn-on surges; the letter gives the rating, not the voltage to design for.
- **Makers' own marks** — date codes, logos, ID codes — share the same face and are not read.
- A value typed in is encoded with R below 10 pF; tantalum parts never go that low.
- **Standard value.** A tolerance letter picks the one series the value is checked against (K,  $\pm 10\%$ , E12; M,  $\pm 20\%$ , E6). A voltage letter says nothing about tolerance, so the value is checked against E6, the series tantalum and polymer capacitors are stocked in.

## What it deliberately does not do

- **The two-character SMD capacitor code** (a letter for the figures and a digit for the multiplier, IEC 60062 Annex B) is not read.
- **Ceramic discs and film parts** have their own code tools.

## How it was checked

Against the AVX datasheet's own examples: 227A reads 220  $\mu\text{F}$  at 10 V, J106 and J 106 read 10  $\mu\text{F}$  at 6.3 V. Also 475V  $\rightarrow$  4.7  $\mu\text{F}$  at 35 V, 106K  $\rightarrow$  10  $\mu\text{F}$   $\pm 10\%$ , 226M  $\rightarrow$  22  $\mu\text{F}$   $\pm 20\%$ .

The review made the tool read a voltage letter printed before the digits (J106 was rejected), applied the 8 and 9 multipliers and R-code length fix from the ceramic tool, limited the units to pF, nF and  $\mu$ F, and made the drawing show the marking as typed.

The re-walk checked the value against the tolerance's own series (156A is an E6 value, 825K an E12 one), showed a value read from a code in its own unit (471 as 470 pF) and put a space in "10 V".

# Series/parallel (C)

## What it computes

The total capacitance of two to four capacitors in parallel or in series, the range it can take given their tolerance, and the nearest single standard value. Each capacitor has a slider that steps through standard values.

## Source

The charge relation  $Q = C \times V$  applied to parts sharing a voltage (parallel) or a charge (series); standard values from IEC 60063.

## What it means

Capacitors combine the **opposite way to resistors**.

**In parallel** — side by side between the same two points. Each has the same voltage across it and stores its own charge, so the charges add, and so do the capacitances — it is as if the plates were one larger plate:

$$C = C1 + C2 + C3 + \dots$$

A 100 nF and a 10 μF in parallel make 10.1 μF. This is how a board gets both a large capacitance for slow supply dips and a small one for fast noise: in parallel they add, and each does the part of the job it is good at.

**In series** — end to end. The same charge has to appear on every capacitor in turn, and their voltages add up to the total, so it is the reciprocals that add:

$$\begin{aligned} 1/C &= 1/C1 + 1/C2 + 1/C3 + \dots \\ C &= C1 \times C2 / (C1 + C2) \quad \text{for two} \end{aligned}$$

The total is always less than the smallest one. Two equal capacitors give half of one; 1 μF in series with 10 μF gives 909 nF, a little less than the 1 μF.

## Voltage in series

The reason to put capacitors in series is usually voltage: two 2.7 V supercapacitors in series make a part that can be charged to 5.4 V. But the voltage does not split evenly on its own. With the same charge on each,

$$V1 / V2 = C2 / C1$$



so the **smaller** capacitor takes the **larger** share of the voltage. Real capacitors differ by their tolerance and by their leakage, so one of them can be pushed past its rating. Series capacitors that are meant to share a voltage — supercapacitors, electrolytics in a high-voltage supply — get a resistor across each one, or an active balancing circuit, to hold the split even.

## Tolerance of the total

As for resistors, if every part is off by the same fraction in the same direction the total is off by exactly that fraction, so the worst-case range is the total  $\times (1 \pm t)$ . Many capacitors have wide tolerances —  $\pm 10\%$  or  $\pm 20\%$  for ceramics and electrolytics, and some ceramics  $+80/-20\%$  — so the range is often the more important number.

## Assumptions and limits

- **Up to four capacitors**, all in series or all in parallel.
- **Nominal values.** Class 2 ceramics lose capacitance with applied DC voltage and with temperature; electrolytics lose it as they age. The total is only as good as the values typed in.
- Leakage, ESR and inductance are ignored.

## What it deliberately does not do

- **Balancing resistor values** for series capacitors are not calculated.
- **Resistors** have their own series/parallel tool.

## How it was checked

By hand: 100 nF + 100 nF in parallel is 200 nF, in series 50 nF; 1  $\mu$ F and 10  $\mu$ F in series give 909.1 nF; two 2.7 F supercapacitors in series give 1.35 F. The ranges are the total  $\times (1 \pm t)$ .

# RC charge and discharge

## What it computes

For a capacitor charging from a supply through a resistor, or discharging through one: the time constant  $\tau$ , the voltage after a given time, or the time it takes to reach a given voltage. A table gives the percentage reached after one to five time constants.

## Source

The solution of the circuit's differential equation, from Ohm's law and the capacitor law  $I = C \times dV/dt$ . No standard is involved.

## What it means

A capacitor stores charge, and its voltage rises as the charge builds up. When it charges through a resistor, the current is largest at the start, when the capacitor is empty and the whole supply voltage is across the resistor. As the capacitor fills, less voltage is left across the resistor, the current falls, and the charging slows. So the voltage does not climb in a straight line but bends over, approaching the supply ever more slowly.


The pace is set by one number, the **time constant**:

$$\tau = R \times C$$

10 k $\Omega$  and 100  $\mu$ F make  $\tau = 1$  second. In each time constant the capacitor covers 63% of the distance that is still left:

| After    | Charging, % of Vs | Discharging, % of V0 left |
|----------|-------------------|---------------------------|
| 1 $\tau$ | 63.2%             | 36.8%                     |
| 2 $\tau$ | 86.5%             | 13.5%                     |
| 3 $\tau$ | 95.0%             | 5.0%                      |
| 4 $\tau$ | 98.2%             | 1.8%                      |
| 5 $\tau$ | 99.3%             | 0.7%                      |

It never quite arrives: mathematically the curve only approaches Vs. In practice **5 $\tau$**  counts as fully charged or discharged.



**Discharging** is the same curve upside down: a charged capacitor connected across a resistor, with no supply, loses 63% of what is left in each  $\tau$ .

## The formulas

|              |                                         |                                   |
|--------------|-----------------------------------------|-----------------------------------|
| Charging:    | $V(t) = V_s \times (1 - e^{(-t/\tau)})$ | $t = -\tau \times \ln(1 - V/V_s)$ |
| Discharging: | $V(t) = V_0 \times e^{(-t/\tau)}$       | $t = -\tau \times \ln(V/V_0)$     |

The second form of each answers the practical question — how long until the capacitor reaches a threshold. For example, how long a power-on reset holds, or when a slowly rising input crosses a logic level: charging to 2/3 of  $V_s$  takes  $\tau \times \ln 3 = 1.1 \tau$ , which is where the 555 timer's 1.1 R C comes from.

## Assumptions and limits

- **An ideal source and capacitor.** The supply is taken as having no internal resistance, and the capacitor as having no leakage. A leaky capacitor, such as an electrolytic, levels off below  $V_s$ .
- **Nominal R and C.** With a  $\pm 20\%$  capacitor the times are  $\pm 20\%$  too.
- The target voltage must lie between 0 and  $V_s$ : a charging capacitor never reaches  $V_s$ , so asking for exactly  $V_s$  has no finite answer.

## What it deliberately does not do

- **Switch debouncing and RC timing for logic inputs** have the Debounce / RC timing tool, which adds input thresholds.
- **Energy** stored in the capacitor is the Stored energy tool.

## How it was checked

By hand: 10 k $\Omega$  and 100  $\mu$ F give  $\tau = 1$  s; at 1 s a 5 V charge reaches 3.161 V (63.21%); 4.5 V is reached at 2.303 s ( $\ln 10 \times \tau$ ); discharging from 5 V, 1.839 V is left at 1 s, and 0.5 V is reached at 2.303 s.

The review fixed the layout, which put the voltage result 24 px under the tab bar on the target phone; R and C now share a line, the voltage and time fields share the next, and the results sit in one row. The schematic no longer shows a supply when discharging, and the resistance units are  $\Omega$ , k $\Omega$  and M $\Omega$ .

# Capacitor stored energy

## What it computes

The energy stored in a charged capacitor from its capacitance and voltage, or either of those from the other two; and the charge it holds.

## Source


The energy of a capacitor charged from zero: the work done moving charge onto it against its own rising voltage. A standard result; no standard document is involved.

## What it means

A capacitor stores energy in the electric field between its plates. Charging it takes work, because each bit of charge has to be pushed on against the voltage already there; that work is held, and comes back out when the capacitor discharges.

$$E = \frac{1}{2} \times C \times V^2 \quad \text{in joules, with } C \text{ in farads and } V \text{ in volts}$$

$$Q = C \times V \quad \text{the charge, in coulombs}$$



**The square matters.** Twice the voltage stores four times the energy. A 1000 µF capacitor at 12 V holds 72 mJ — a small spark; the same capacitor at 400 V would hold 80 J. This is why the capacitors in mains power supplies, camera flashes and microwave ovens are dangerous long after the power is removed, and why they are discharged through a resistor before being touched.

**How much is a joule.** One joule is one watt for one second. A 1 F supercapacitor at 2.7 V holds 3.6 J: enough to keep a 10 mW real-time clock circuit running for six minutes.

Rearranged:

$$V = \sqrt{2E / C} \quad \text{the voltage needed to store } E \text{ in } C$$

$$C = 2E / V^2 \quad \text{the capacitance needed to store } E \text{ at } V$$

## Usable energy

A circuit seldom uses all of it. A device powered from a capacitor stops working when the voltage falls below its minimum, so the energy it can draw is the difference between two voltages:

$$E_{\text{usable}} = \frac{1}{2} \times C \times (V_{\text{start}}^2 - V_{\text{min}}^2)$$

A supercapacitor charged to 5 V feeding a circuit that needs at least 3 V delivers only  $1 - (3/5)^2 = 64\%$  of what it holds. Use the tool twice, at the two voltages, and subtract.

## Assumptions and limits

- **An ideal capacitor.** Leakage slowly drains the stored energy; ESR turns some of it to heat when it is drawn fast.
- **Nominal capacitance.** Class 2 ceramics hold less than their marked value at high voltage, and electrolytics lose capacitance with age.

## What it deliberately does not do

- **Discharge time** through a load is the RC charge and discharge tool's job.
- **Hold-up time** from the usable energy is not calculated directly; it is the usable energy divided by the power drawn.

## How it was checked

By hand: 1000  $\mu\text{F}$  at 12 V holds 72 mJ and 12 mC; 1 F at 2.7 V holds 3.645 J; storing 1 J in 100  $\mu\text{F}$  needs 141.4 V; storing 10 J at 400 V needs 125  $\mu\text{F}$ .


# Inductor color code

## What it computes

The inductance and tolerance that the colour bands on an axial (leaded) inductor stand for, and the other way round: the bands for a value. A second pill handles the 5-band form with the wide silver identifier band.

## Source

- **Inductors Inc., colour band guide:** the colours carry the same digits and multipliers as the resistor code, read in microhenries; gold among the digit bands is the decimal point; the tolerance colours are gold  $\pm 5\%$ , silver  $\pm 10\%$  and black  $\pm 20\%$ ; the wide silver band is a "military identifier" that does not imply military qualification.
- **Vishay IM series** (MIL-PRF-15305 molded inductors): inductance tolerances of  $\pm 1$ ,  $\pm 3$ ,  $\pm 5$ ,  $\pm 10$  and  $\pm 20\%$ , of which only the last three have a colour.
- IEC 60062, the standard for resistor and capacitor codes, does not cover inductors; the inductor code borrows its colours.



## What the bands mean

Small axial inductors look like fat resistors, and use the same colours — but the value is in **microhenries (µH)**, not ohms.

- **Band 1 and 2** — the two digits (black 0, brown 1, red 2 ... white 9).
- **Band 3** — the multiplier, how many zeros follow: black  $\times 1$ , brown  $\times 10$ , red  $\times 100$ ; gold  $\times 0.1$  and silver  $\times 0.01$  for small values.
- **Band 4** — the tolerance: gold  $\pm 5\%$ , silver  $\pm 10\%$ , black or no band  $\pm 20\%$ .

So brown–black–black–gold is  $10 \times 1 = 10 \mu\text{H} \pm 5\%$ , and red–red–brown–silver is  $22 \times 10 = 220 \mu\text{H} \pm 10\%$ .

**Gold as the decimal point.** A gold band in first or second place is not a digit but the decimal point, and the band after it is then a digit too: yellow–gold–violet is  $4.7 \mu\text{H}$ , gold–yellow–violet is  $0.47 \mu\text{H}$ . The same  $4.7 \mu\text{H}$  can also be written yellow–violet–gold, with gold as a  $\times 0.1$  multiplier; both forms are read.

**The 5-band form.** Some parts start with a wide silver band. It is an identifier in the military style and carries no digit; the four bands after it read as above. It does not by itself mean the part is military-qualified.

## Assumptions and limits

- **Tolerance colours are only three.** Precision inductors ( $\pm 1\%$ ,  $\pm 2\%$ ,  $\pm 3\%$ ) are not marked with a tolerance colour; their tolerance is on the datasheet or in the part number.

- **Makers vary.** Colour-coded inductors are old enough, and made by enough small makers, that some use their own scheme. When the result looks implausible for the part's size, measure it.
- Inductance alone does not say what current the part can carry; that is on its datasheet.

## What it deliberately does not do

- **SMD inductor markings** are the SMD code tool.
- **Resistor bands** are the resistor Color code tool, which uses the full IEC 60062 tolerance colours.

## How it was checked

Brown–black–black–gold reads 10  $\mu\text{H}$   $\pm 5\%$ ; red–red–brown–silver reads 220  $\mu\text{H}$   $\pm 10\%$ ; yellow–gold–violet 4.7  $\mu\text{H}$ ; gold–yellow–violet 0.47  $\mu\text{H}$ ; typing 4.7  $\mu\text{H}$  gives yellow–violet–gold.

The review found the tolerance list offering the resistor's precision colours, brown  $\pm 1\%$  down to grey  $\pm 0.01\%$ , which no colour-coded inductor uses; it is now gold, silver, black and none. It added gold as a decimal point in the digit bands, which the tool could not read, and corrected the note's claim that the silver band marks a part as military-qualified.

# Karnaugh map simplification

## What it computes

Takes a truth table for 2, 3 or 4 variables — each cell either 0, 1 or a don't-care — and returns the minimum sum-of-products expression, with each product term drawn as a group on the map in its own colour.

## Source

The map is Karnaugh's (1953), itself a reworking of the Veitch diagram. The minimisation is Quine–McCluskey, which is an exact algorithm rather than a heuristic; the covering step is solved by exhaustive search, which is affordable at four variables and would not be at twenty.

## Why the formulas are these

The whole method rests on the ordering along the edges: **00, 01, 11, 10**, not 00, 01, 10, 11. That is Gray code, and it means any two side-by-side cells differ in exactly one variable. When two adjacent cells are both 1, that variable appears once true and once complemented, so it cancels:

$$A \cdot \bar{B} \cdot C + A \cdot B \cdot C = A \cdot C \cdot (\bar{B} + B) = A \cdot C$$

Extend it and a rectangle of  $2^k$  cells cancels  $k$  variables. A pair drops one, a block of four drops two, eight drops three. That is the entire trick — the map is just an arrangement that puts algebraically adjacent terms physically next to each other.

**The edges wrap.** The first and last columns also differ in one variable (00 and 10 differ only in the middle bit), and so do the first and last rows. The map is really a torus. That is why the four corners of a four-variable map form one group of four, and why a group can run off one edge and continue on the other. A torus rectangle has no single box on a flat sheet, so the tool draws wrapping groups as two or four pieces in the same colour.

**Quine–McCluskey.** Terms are combined pairwise wherever they differ in one bit, the shared bit becoming a dash; repeat until nothing more combines. What never combined is a **prime implicant** — a group that cannot be made larger. Then a cover must be chosen: a minterm covered by exactly one prime implicant forces that implicant in (it is *essential*), and the rest is a set-cover problem solved here by branch and bound.

The cover is ordered **fewest terms first, then fewest literals**. The second key matters: several covers can tie on term count while differing in how many variables survive, and the answer a textbook prints is the one with fewer letters.

**Don't-cares** are included when forming groups but never required to be covered. That is where most of the saving usually comes from: an X adjacent to a group lets the group double in size, dropping another variable



for free.

**Parity is what the map cannot see.** A checkerboard of ones has no two adjacent cells, so every group is a single cell and the sum of products comes out at full length — twelve of twelve literals on a three-variable map, nothing saved at all. That function is exclusive-OR:

$$A \oplus B \oplus C = 1 \text{ whenever an odd number of } A, B, C \text{ are } 1$$

Two XOR gates implement it against four ANDs and an OR for the sum of products. The map cannot find this because XOR is not a sum of products in any small form — its minimal SOP genuinely is the long one. The tool therefore tests, separately from the map, whether the function is the parity of any subset of its variables (or the complement of one) and prints that form when it is.

## Assumptions and limits

- Two to four variables. Five and six variable maps exist, drawn as stacked pairs, but they stop being readable and the whole point of the map is that you can see it.
- The result is a minimum, not *the* minimum: several different expressions can tie on both term and literal count, and the tool shows one of them.
- Literal count is a proxy for cost, not a gate count. It ignores that inverters may be free if a signal is already available complemented, and that fan-in is limited in real logic families.
- No account of hazards. A minimal cover can contain a static hazard — a glitch when two inputs change at once — which is removed by adding a redundant group that the minimiser will never choose because it costs a term.

## What it deliberately does not do

- **Product of sums.** The minimal POS is the same procedure applied to the zeros, then inverted. It would double the interface for a form that is asked for far less often.
- **More than four variables,** as above.
- **Hazard-free covers.** Adding the redundant terms is a different objective from minimisation and would make the answers disagree with what a textbook shows.

## How it was checked

A second minimiser was written independently for the purpose: enumerate every possible cube, discard those covering a zero, and find the smallest cover by iterative deepening. The two were compared over 3500 random functions at 2, 3 and 4 variables, and 12000 more were checked only for the expression reproducing the truth table at every care position.

That comparison earned its keep. Term counts matched from the start, but 34 of the first 3400 came back with more literals than the reference — the cover search was returning whichever minimum-size cover it reached first, with no tie-break on literals. A correct minimum by term count, and still not the printed answer.

The parity detector was checked exhaustively over all 256 three-variable functions: exactly eight are parity forms (three pairs and one triple, each with its complement), and each one's exclusive-OR form reproduces its truth table.

# I2C pull-up resistor

## What it computes

The window of pull-up resistances that work on an I<sup>2</sup>C bus, given the supply, the bus capacitance and the speed mode — and a suggested value inside it, snapped to a standard E-series.

## Source

The I<sup>2</sup>C specification, **NXP UM10204**. Every constant here is from it: the 0.4 V output-low level, the 3 mA and 20 mA sink currents, the rise-time budgets, and the bus capacitance ceilings.

## Why the formulas are these

I<sup>2</sup>C lines are **open-drain**. Nothing on the bus drives a line high; devices can only pull it low, and the pull-up resistor is what brings it back up. That single fact produces both bounds, pulling in opposite directions.

**The lower bound is the pull-down.** When a device asserts the line, its open-drain transistor has to hold it below the 0.4 V that counts as a low, and the current it must swallow is whatever the pull-up pushes through:

$$R_{p \text{ min}} = (V_{dd} - 0.4 \text{ V}) / I_{OL}$$

The spec requires devices to sink **3 mA** up to Fast mode, and **20 mA** in Fast mode Plus. A smaller resistor asks for more than the transistor is rated to take, and the line stops reaching a valid low.

**The upper bound is the rise.** Going up, the line is just an RC: the pull-up charging the bus capacitance. The spec measures rise time between **0.3 V<sub>dd</sub>** and **0.7 V<sub>dd</sub>**, and an RC charge reaches those fractions at


$$\begin{aligned} t(0.3 \text{ V}_{dd}) &= RC \cdot \ln(1/0.7) \\ t(0.7 \text{ V}_{dd}) &= RC \cdot \ln(1/0.3) \end{aligned}$$

so the specified rise time is the difference:

$$t_r = RC \cdot [\ln(1/0.3) - \ln(1/0.7)] = RC \cdot \ln(7/3) = 0.8473 \cdot RC$$

Hence

$$R_{p \text{ max}} = t_r / (0.8473 \times C_b)$$



That 0.8473 is copied everywhere without explanation; it is just  $\ln(7/3)$ . The budget for  $t_r$  is 1000 ns at 100 kHz, 300 ns at 400 kHz and 120 ns at 1 MHz.

**The suggestion** is the geometric mean of the two bounds,  $\sqrt{R_{p \min} \times R_{p \max}}$ , snapped to the chosen E-series. The geometric mean is the centre of the window on a logarithmic scale, which is what "equal margin on both sides" means when the two ends can differ by a decade. For a 3.3 V standard-mode bus with 100 pF it lands on 3.3 k $\Omega$  — the value most designs actually fit.

**When the window closes.**  $R_{p \max}$  falls as bus capacitance rises, while  $R_{p \min}$  does not move. Past roughly 366 pF in Fast mode at 3.3 V the two cross and no resistor satisfies both: the bus cannot rise fast enough without asking the pull-down for more current than it is rated for. That is a real dead end, not a poor choice, and the ways out are a shorter bus, fewer devices, a slower mode, or an active bus buffer that isolates segments.

## Assumptions and limits

- **Bus capacitance is an estimate.** Reckon on a few pF per centimetre of track plus about 10 pF per device. It is the input people are most wrong about, and the whole upper bound scales with it. Measure it on a long bus.
- The RC model treats the bus as one lumped capacitance and the pull-up as ideal. At Fast mode Plus, with long tracks, transmission-line behaviour starts to matter and the model flatters the real rise time.
- Leakage from many devices and from the pull-up itself is ignored. On a bus with dozens of devices it adds a small standing current that eats into the low-level margin.
- $I_{OL}$  defaults to the spec value for the mode. A specific part may be rated differently; the field is editable for that reason.

## What it deliberately does not do

- **Current-source or active pull-ups.** Fast mode Plus devices sometimes use a current source rather than a resistor, which removes the RC ceiling entirely. Different circuit, different arithmetic.
- **Bus buffers and repeaters.** They break a bus into segments each with their own capacitance and their own resistor; the tool sizes one segment, so run it once per segment.
- **Level shifters.** A MOSFET level shifter between two voltage domains puts a pull-up on each side, and the two interact. Not modelled.

# UART baud rate

## What it computes

The divisor a UART needs to reach a wanted baud rate from a given peripheral clock, the rate it actually produces once that divisor is rounded to an integer, and the resulting error — with the error shown against every standard rate on the same clock.

## Source

There is no single standard for baud generation: it is per-device, and the register is called BRR, UBRR, DLL/DLM or SBRG depending on whose part it is. What is common to nearly all of them is the shape — an integer divider off a peripheral clock, with the receiver oversampling each bit — and 16x oversampling, inherited from the 8250 and 16550 and still the default almost everywhere. The error budget below is derived from the sampling behaviour, not quoted.

## Why the formulas are these

A UART has no oscillator of its own. It divides whatever clock it is given:

$$\text{Divisor} = \text{Clock} / (\text{Oversample} \times \text{Baud})$$

The divider is a counter, so the value written is an **integer**. Rounding it is the whole problem:


$$\text{Actual baud} = \text{Clock} / (\text{Oversample} \times \text{round}(\text{Divisor}))$$

At 16 MHz, 115200 baud and 16x oversampling the divisor comes to 8.68. Written as 9, the line runs at 111111 baud — 3.55% slow. That single fact is why 14.7456 MHz crystals exist:  $14745600 / 16 = 921600$ , which divides exactly into every standard rate, so the divisor is always a whole number and the error is zero. 11.0592 MHz and 18.432 MHz are the same trick at other speeds.

**Why the error matters, and how much is allowed.** There is no clock shared between the two ends. The receiver finds the falling edge of the start bit, starts its own counter, and samples each bit at what it believes is the middle. Take the receiver as the reference: it samples bit  $n$  at time

$$(n + 0.5) \times T_{rx}$$

while the transmitter's bit  $n$  occupies



$$[ n \times T_{tx} , (n+1) \times T_{tx} ]$$

Writing  $r = T_{tx} / T_{rx}$ , the sample lands in the right bit only while

$$n \times r < n + 0.5 < (n+1) \times r$$

Nothing resynchronises inside a frame, so the worst case is the last bit. For 8N1 that is the stop bit,  $n = 9$ :

$$\begin{aligned} 9r < 9.5 & \rightarrow r < 1.0556 \\ 9.5 < 10r & \rightarrow r > 0.95 \end{aligned}$$

So the two clocks may differ by about **−5.0% to +5.56%** — usually quoted as  $\pm 5\%$ , and slightly asymmetric because the sample point is fixed at mid-bit while the bit it must land in is what stretches. That budget covers **both** ends together, which is why 2% each is the practical rule: it leaves room for the other end plus edge jitter and finite rise time.

Note what does *not* happen: the error does not accumulate across a message. The receiver re-finds the start edge on every frame, so the drift resets ten bits at a time. A 3% error is not 3% worse each byte; it is the same 3% on each of them, which is why it either works or does not.

**The drawing shows exactly this.** The waveform is drawn at the transmitter's real bit width, stretched or squeezed by the error, while the sample marks stay evenly spaced where the receiver puts them. Raise the error and the marks visibly walk toward the bit edges; the ones that fall outside their own bit turn red. It is the inequality above, drawn.

## Assumptions and limits

- **8N1 is assumed** — one start bit, eight data, one stop, no parity. Adding parity or a second stop bit makes the frame longer and the budget tighter, since the worst case moves to a later bit.
- **The receiver is taken as exact.** In reality both ends have error and the two add; the tool reports one side's, which is why the caution appears at 2% rather than 5%.
- **Integer divisors only.** See below.
- Oversampling is a divider stage, not noise immunity: 8x doubles the top speed from a given clock and halves the margin around each sample, so a noisy line suffers more.

## What it deliberately does not do

- **Fractional dividers.** Many modern UARTs — STM32, SAM, nRF — add fractional bits below the integer divisor, which cuts the error by roughly the resolution of that fraction and would make the numbers here pessimistic for those parts. Modelling it means picking a vendor, since the fraction is 4 bits on some parts, 6 on others, and a full 16-bit accumulator on a few.
- **Auto-baud detection**, where the receiver measures a known character instead of being told the rate.
- **Anything above the physical layer:** framing errors, break detection, flow control.

## How it was checked

The classic cases were recomputed by hand: 16 MHz at 115200 gives divisor 8.68, used 9, 111111 baud, -3.549%; 14.7456 MHz gives exactly 8 and 0%; 16 MHz at 9600 gives 104.17, used 104, +0.160%, which is why that part runs 9600 happily and 115200 badly. 8 MHz at 115200 gives +8.507%.

The sample-point test in the drawing was checked against the inequality above at several errors: at 0% and -3.5% every sample lands inside its bit, at +0.16% likewise, and at +8.5% the last four fail — which is what  $n \times 0.9216 < n + 0.5$  predicts, the first failure being at  $n = 6$ .

# Crystal load capacitance

## What it computes

The two capacitors a Pierce oscillator needs so the crystal sees the load it was cut for, the load a given pair actually presents, and how far off frequency the difference pulls it.

## Source

The series expression appears in every oscillator application note — ST AN2867, the Abracon and Epson design guides, and the NXP equivalents. The pulling expression comes from the crystal's equivalent circuit, the Butterworth–Van Dyke model: a motional arm of L1, C1 and R1 in series, all in parallel with the shunt capacitance C0 of the electrodes and holder.

## Why the formulas are these

**The two capacitors are in series, not in parallel.** This is the step people get backwards. Each one runs from one crystal terminal to ground, so the loop the crystal drives is: out through CL1, along the ground, back through CL2. Two capacitors in series:

$$CL1 \times CL2 / (CL1 + CL2)$$

which for an equal pair is just half of one of them. Add the stray capacitance of the two tracks and the two oscillator pins — that sits directly across the same nodes, so it adds on:


$$\text{Load} = CL1 \times CL2 / (CL1 + CL2) + \text{Stray}$$

Turn it round for a symmetric pair and the capacitors come out at

$$CL1 = CL2 = 2 \times (CL \text{ spec} - \text{Stray})$$

A crystal marked 12.5 pF on a board with 3 pF of stray therefore wants **19 pF** parts, not 25 pF and not 12.5 pF. Both of those are common mistakes, and the second is the worse one.

**Why the load changes the frequency at all.** A crystal has two resonances. At series resonance the motional arm is purely resistive; a little above it the arm looks inductive and resonates with everything capacitive across it — C0 and whatever load the circuit adds. The parallel-resonant frequency is



$$f_p = f_s \times [ 1 + C_1 / (2 \times (C_0 + C_L)) ]$$

A crystal is cut and trimmed so that this lands on the marked frequency **for one particular CL**. Present a different load and it lands somewhere else, which is the whole reason the number is on the datasheet.

Differentiating that with respect to CL gives the sensitivity, which is what the tool reports as pullability:

$$d(\Delta f/f) / dC_L = - C_1 / (2 \times (C_0 + C_L)^2)$$

The sign is negative: **more load, lower frequency**. For a typical MHz part with  $C_1 = 8 \text{ fF}$  and  $C_0 = 3 \text{ pF}$  at  $C_L = 12.5 \text{ pF}$  that comes to about 16.6 ppm/pF, so fitting 22 pF where 19 pF was wanted — 1.5 pF too much load — pulls it roughly 25 ppm slow. A 32.768 kHz tuning fork has a tenth the motional capacitance and pulls far less per pF, around 6 ppm/pF on the same load, which is why watch crystals are specified so tightly: there is little room to trim them back.

## Where the inputs come from, and what they are called

Three of the inputs are easy to look for in the wrong place.

- **C0** — on the crystal datasheet, as  $C_0$ , "shunt capacitance" or "static capacitance". Often given only as a maximum; that is usable.
- **C1** — the motional capacitance,  $C_1$  on datasheets that give it, after the Butterworth–Van Dyke model. **Many do not.** When it is missing, look for a pullability figure in ppm/pF, or a "frequency versus load capacitance" curve, and type that into Pullability instead: the tool works C1 back from it. C1 is the one that stays fixed when the load changes — it is the crystal — while pullability is its slope at whatever load is chosen, so changing CL spec recomputes pullability, not C1.
- **Stray** — **never** on the crystal datasheet, because it is not the crystal. It is the capacitance of the two oscillator pins, sometimes listed in the MCU datasheet, plus the two tracks. Application notes such as ST AN2867 give 2–5 pF as the usual range; that is an order of magnitude, not a value to copy.

The load capacitors are called **CL1 and CL2** rather than C1 and C2, which is how MCU application notes label them and which keeps C1 free for its datasheet meaning. Every capacitance here only ever comes in one unit — pF for load, stray and shunt, fF for motional — so the fields show the unit fixed rather than offering a picker.

**When CL says "Series"**. Some datasheets offer, or list, a series-resonant cut: the crystal is trimmed to its marked frequency with no load at all. A Pierce oscillator cannot present no load, so the part runs above its marking by the whole parallel-resonance offset,

$$\Delta f/f = C_1 / (2 \times (C_0 + \text{Load}))$$

which is not a trim but a few hundred ppm: 8 fF against 3 pF + 12.5 pF is +258 ppm. Type 0 in CL spec to see it. The fix is to order the part with a load capacitance, or to use it in a circuit that runs it at series resonance.

## Assumptions and limits

- **Stray is an estimate and it dominates the mistakes.** Two to five picofarads is the usual range for the pins plus short tracks, but a long or guarded layout can be well outside it. If the frequency matters, measure.
- The presets are typical for each family, not for any specific part, and a wrong C1 scales the pullability proportionally.
- The pulling figure is a **slope taken at the specified load**, so it is exact for small errors and increasingly optimistic for large ones, since the true curve flattens as CL grows.
- Sizing assumes a **symmetric pair**. Asymmetric capacitors are legitimate and the tool will report the load they present, but it will not propose them.

## What it deliberately does not do

- **Drive level and negative-resistance margin.** Whether the oscillator starts and whether it overdrives the crystal are the other half of the design, set by the amplifier's transconductance, the crystal's ESR and any series resistor. It is a separate calculation with separate datasheet numbers.
- **The feedback resistor**, which is inside the MCU on every part this applies to and is not drawn for that reason.
- **Temperature and ageing.** The load error here is a fixed offset; drift over temperature is the crystal's cut and belongs with the ppm tooling.

## How it was checked

Recomputed by hand: 12.5 pF specified with 3 pF stray gives 19 pF capacitors and a presented load of exactly 12.5 pF, 0 ppm. Fitting 22 pF instead presents 14 pF, 1.5 pF over, and pulls  $-24.97$  ppm against  $16.65 \times 1.5 = 24.98$  by hand. Fitting 33 pF presents 19.5 pF and pulls  $-116.5$  ppm against  $16.65 \times 7 = 116.6$ . The watch preset returns 6.378 ppm/pF against  $2.5 \text{ fF} / (2 \times (1.5 + 12.5 \text{ pF})^2) = 6.38$  by hand.

The pullability was wrong by twelve orders of magnitude on first write:  $C1 / (2(C0+CL)^2)$  is *per farad*, not dimensionless, because the numerator is first order in capacitance and the denominator is second. It needs multiplying by one picofarad to become per-pF. The hand check is what caught it.

# Oscillator stability

## What it computes

The total frequency error budget of a crystal oscillator — tolerance, temperature, ageing and load pulling combined — as a worst case and as a root-sum-square estimate, then expressed as hertz and as time gained or lost per day and per year.

For a single ppm figure applied to a single value, the **PPM converter** under Tools does that; this tool is the budget that produces the figure.

## Source

The contributions and their names are the ones on crystal datasheets: frequency tolerance at a reference temperature, frequency stability over the operating temperature range, and ageing, usually per first year. The reference is 20 °C on some datasheets and 25 °C on others — one Petermann-Technik page quotes both for the same part — so the field is labelled Tolerance without a temperature. The 32.768 kHz temperature curve is the parabola the tuning-fork datasheets from Epson, Abracon and Micro Crystal all quote, with a coefficient of about  $-0.034 \text{ ppm/}^{\circ}\text{C}^2$  ( $\pm 0.006$ ) about a turnover near 25 °C.

## Why the formulas are these


**Two kinds of error, kept apart.** Tolerance, temperature stability and ageing are *limits*: the part is somewhere within  $\pm$  each of them, and which way it sits is not known. The tuning fork's temperature drift and the load pulling are *offsets*: they have a known sign and they move the whole window. Mixing the two kinds is how budgets end up either too pessimistic or wrong-signed, so the tool computes

$$\begin{aligned}\text{Offset} &= \text{Drift} + \text{Pulling} \\ \text{Worst case} &= \text{Offset} \pm (\text{Tolerance} + \text{Temperature} + \text{Ageing})\end{aligned}$$

and reports the extreme of that window with its sign. A window of  $-69$  to  $-23$  ppm is a clock that can only lose time, and writing it as  $\pm 69$  ppm would throw that away.

**Worst case and RSS.** The worst case adds every limit at full size in the same direction. It is what datasheets call *overall stability*, and it is the right figure to design against when failure is not an option. If the contributions are independent, all of them sitting at their extremes together is unlikely, and the root-sum-square

$$\text{RSS} = \text{Offset} \pm \sqrt{(\text{Tolerance}^2 + \text{Temperature}^2 + \text{Ageing}^2)}$$



is the likelier spread. Offsets are added directly, never squared: they are not random.

**The tuning-fork parabola.** A 32.768 kHz crystal is a quartz tuning fork, and its frequency falls off either side of a turnover temperature:

$$\Delta f/f = -0.034 \text{ ppm/}^{\circ}\text{C}^2 \times (T - 25 \text{ }^{\circ}\text{C})^2$$

It only ever goes down. At 0 °C that is –21 ppm, about 1.8 seconds a day; at –20 °C it is –69 ppm, six seconds a day. That is why a watch or a data logger left outdoors in winter runs slow, and why the tool asks for the temperature the fork will actually see rather than a spec. MHz crystals are AT-cut plates, whose curve is a cubic set by the cut angle; their datasheets fold it into one ±ppm figure over the range, which is what the MHz pill asks for.

**Into time.** A day is 86 400 s, so 1 ppm is 86.4 ms a day, and a year of 31 557 600 s makes 1 ppm about 31.6 s a year. Twenty ppm is 1.73 s a day or 10.5 minutes a year.

## Assumptions and limits

- **Ageing is not linear.** It is fastest in the first months and slows roughly logarithmically, so ten years is much less than ten times the first-year figure. Enter the total you expect over the period, not a rate.
- **The parabola is typical,** not the part's own curve: the coefficient varies by about ±18% and the turnover by several degrees. Near 25 °C that barely matters; at the extremes it does.
- The MHz temperature figure is taken as a symmetric ± limit, which is how it is specified, even though the real AT-cut curve is not symmetric.
- Pulling is carried in as a signed figure from the crystal load tool and is treated as exact.
- **Tolerance never includes temperature on a crystal datasheet.** The temperature stability is a separate line, often a separate ordering table, and the two add: ±20 ppm tolerance with ±30 ppm over –40/+85 °C is up to ±50 ppm before ageing. A packaged **XO or TCXO** is different: its "frequency stability" is often one all-inclusive figure covering tolerance, temperature, supply and load, sometimes ageing too. Read the footnote; if it is all-inclusive, enter it as Tolerance and leave the rest at zero.

## What it deliberately does not do

- **Compensated oscillators.** A TCXO or an RTC with digital trimming cancels most of the temperature term; its datasheet gives a residual stability, which can be entered as Temp. stability on the MHz pill.
- **Short-term stability and jitter.** Phase noise and Allan deviation are a different question from where the average frequency sits.
- **Protocol limits.** Whether a budget meets USB, CAN or Ethernet timing needs those standards' own tolerances and belongs with the tools for each.

## How it was checked

By hand, MHz defaults: 20 + 30 + 3 = 53 ppm worst case;  $\sqrt{(400 + 900 + 9)} = 36.18$  ppm RSS; 53 ppm of 16 MHz is 848 Hz; 53 × 86.4 ms = 4.579 s a day; 53 × 31.56 s = 27.88 minutes a year. With –25 ppm of pulling the window becomes –78 to +28 ppm and the worst case reads –78, signed.

32.768 kHz at 0 °C: drift  $-0.034 \times 25^2 = -21.25$  ppm; with  $\pm 23$  ppm of limits the window is  $-44.25$  to  $+1.75$  ppm, worst  $-44.25$  ppm,  $-3.82$  s a day. Switching pills resets pulling, which on first build carried over from one family to the other.

# PLL multiplication factor


## What it computes

Every combination of input divider M, feedback multiplier N and output divider P that a chip's PLL allows, walked exhaustively, with the ones that land closest to a target frequency listed first — each with its error in ppm and the frequencies at the phase detector and the VCO, so the limits can be seen being respected.

## Source

The limits are taken from the vendors' own code rather than from summaries, which disagree with each other:

- **STM32F4** — ST's HAL driver ([stm32f4xx\\_hal\\_rcc.h](#), [stm32f4xx\\_hal\\_rcc\\_ex.h](#)): PLLM 2–63, PLLN 50–432, PLLP 2, 4, 6 or 8, PLLQ 2–15, and the VCO input between 1 and 2 MHz. The VCO output window of 100–432 MHz and the 168 MHz system clock of the F405/407 are from RM0090.
- **RP2040** — the Pico SDK's own calculator, [vcocalc.py](#): reference divider 1–63 with at least 5 MHz after it, FBDIV 16–320, VCO 750–1600 MHz, and two post-dividers each 1–7. The 133 MHz system clock rating is the datasheet's.



## Why the formulas are these

An integer-N PLL has one loop and three dividers:

$$\begin{aligned} \text{PFD} &= \text{Input} / M \\ \text{VCO} &= \text{PFD} \times N \\ \text{Out} &= \text{VCO} / P \end{aligned}$$

The input is divided down to the **phase-frequency detector** (PFD) rate. The VCO runs at N times that, because the loop compares the VCO divided by N against the PFD reference and steers the VCO until they match. The output is the VCO divided down again. Each stage has a window it must stay inside: the PFD because the detector and its filter are designed for a range, the VCO because the oscillator only tunes over a range. That is what makes this a search rather than a division — the obvious M and N often put one of the stages outside its window.

**Ranking.** Rows are ordered by error first. Among equally good rows:

- on the STM32F4, a row whose VCO also divides by some Q to exactly 48 MHz comes first, since USB OTG will not work on anything else;

- then the **higher PFD**. ST's manual recommends a 2 MHz VCO input to limit jitter, and in general a faster comparison lets the loop correct the VCO more often and filter less, which is quieter;
- then the higher VCO.

That is why 8 MHz to 168 MHz on the STM32F4 comes out as M 4, N 168, P 2, Q 7 (PFD 2 MHz, VCO 336 MHz, USB 48 MHz) ahead of the equally exact M 8, N 336, which runs the detector at 1 MHz.

**RP2040's two post-dividers** multiply, so the tool offers each distinct product once, as the pair with the larger first divider — the SDK prefers that split for power.

## Assumptions and limits

- **Integer-N only.** Many modern PLLs add a fractional part to N, which can hit almost any target but trades it for spurs and depends on each vendor's modulator. Not modelled.
- The preset limits are for the named parts. Other members of a family can differ — the F401, for one, needs its VCO in 192–432 MHz, and the RP2350 is not the RP2040 — so use Generic with the right reference manual when in doubt.
- Lock time, loop-filter design and phase noise are not addressed; the tool answers only which settings are legal and how close they land.
- On Generic, the search stops at three million combinations to stay responsive and says so; narrow the ranges for a complete answer.

## What it deliberately does not do

- **Vendor clock trees beyond the PLL** — AHB and APB prescalers, peripheral clock muxes, flash wait states for the chosen frequency. Those are the next steps after this one, and every family arranges them differently.
- **Other presets.** The Generic pill covers any integer-N PLL given its manual; a preset is only worth adding for a part used often enough to save the reading.

## How it was checked

Against the configurations the vendors ship: 8 MHz to 168 MHz on the STM32F4 gives M 4, N 168, P 2, Q 7 with USB at exactly 48 MHz, and from a 25 MHz crystal M 25, N 336, P 2, Q 7 — the standard setting for that crystal. 180 MHz is flagged as above the part's 168 MHz rating and as leaving USB at 45 MHz.

On the RP2040, 12 MHz to 125 MHz gives VCO 1500 MHz with post-dividers 6 and 2, and 133 MHz gives VCO 1596 MHz with 6 and 2 — both exactly the Pico SDK's own default settings. The full RP2040 search takes about 3 ms.

# ADC resolution / quantization

## What it computes

For an ideal N-bit converter and its reference: the size of one step (the LSB), the code a given input produces, the voltage that code stands for, and the quantization error between the two, in volts and in LSB. It works both ways: type a voltage to get the code, or type a code — say, one read off a register — to get the voltage. The range of inputs and codes is shown under the fields.

Two pills cover the two ways converters are built:

- **Single-ended** — 0 to VREF, codes 0 to  $2^N - 1$ . Every MCU ADC (AVR, STM32, RP2040, ESP32) works this way.
- **Bipolar  $\pm$**  —  $-FSR$  to  $+FSR$ , codes  $-2^{N-1}$  to  $2^{N-1} - 1$  in two's complement. Differential converters such as the ADS1115 work this way.

## Source

- The ideal transfer function is the one in IEEE Std 1241 and the one MCU datasheets measure their errors against. The ATmega328P datasheet (ADC characteristics) defines offset error as the deviation of the first transition from its ideal place "at 0.5 LSB", gain error against a last transition 1.5 LSB below full scale, and gives  $\pm 0.5$  LSB as the ideal absolute accuracy.
- The bipolar pill follows the TI ADS1115 datasheet (SBAS444E): the full-scale range is set by the PGA to  $\pm 6.144$ ,  $\pm 4.096$ ,  $\pm 2.048$  (the power-up default),  $\pm 1.024$ ,  $\pm 0.512$  or  $\pm 0.256$  V; Table 7-1 gives the LSB as 125  $\mu$ V at  $\pm 4.096$  V and 62.5  $\mu$ V at  $\pm 2.048$  V; output is binary two's complement, clipping at 7FFFh and 8000h.


## Why the formulas are these

**The step.** N bits give  $2^N$  codes. Spread over the reference:

Single-ended:  $LSB = VREF / 2^N$   
 Bipolar:  $LSB = 2 \times FSR / 2^N$

The bipolar span is twice FSR because it runs from  $-FSR$  to  $+FSR$ . The ADS1115 datasheet writes its Equation 4 as  $LSB = FSR / 2^{16}$ , but its own table (125  $\mu$ V at  $\pm 4.096$  V) only works with FSR as the whole 8.192 V span — the tool asks for the  $\pm$  figure, as the datasheet's register table quotes it, and doubles it.

**The code.** Code k stands for  $k \times LSB$  and covers half an LSB either side, so the ideal converter rounds:



$$\text{Code} = \text{round}(\text{Input} / \text{LSB})$$

The first transition, 0 to 1, is at  $\frac{1}{2}$  LSB. That is why the quantization error is  $\pm\frac{1}{2}$  LSB rather than 0 to  $-1$  LSB: the steps are centred on the ideal straight line, and the error is the sawtooth between the line and the staircase. The drawing shows eight codes around the input at scale, with the dashed ideal line through the step centres.

**The top of the range.** The highest code is  $2^N - 1$ , not  $2^N$ , so the largest voltage a code stands for is  $V_{\text{REF}} - 1$  LSB: 3.2992 V for 12 bits on 3.3 V. An input at  $V_{\text{REF}}$  itself reads the top code, a whole LSB short.

**Two's complement.** A bipolar converter's register holds negative codes as  $2^N + \text{code}$ , so  $-1$  reads FFFFh and  $-\text{FS}$  reads 8000h. The Hex cell shows the register as it would be read, the Code field the signed value it means.

**Dividing by  $2^N - 1$ .** Arduino's ReadAnalogVoltage example computes  $\text{code} \times 5.0 / 1023$ , and ST's LL macro divides by 4095. That treats the top code as exactly  $V_{\text{REF}}$ , stretching the scale: the reading is  $k / (2^N - 1)$  instead of  $k / 2^N$  of  $V_{\text{REF}}$ , high by up to one LSB near the top. Small beside a real ADC's errors, but a systematic one, and it is why a 3.3 V input "reads 3.3 V" on those examples and 3.2992 V here.

## Assumptions and limits

- **An ideal converter.** Offset, gain error, INL, DNL and noise are not modelled; on a real MCU ADC they add up to several LSB, far more than the quantization shown here. Their datasheet figures are in LSB, which is what this tool gives the size of.
- **$V_{\text{REF}}$  is taken as exact.** On most MCUs it is the supply, so the supply's own tolerance scales every reading. A 1% regulator is 41 LSB at 12 bits.
- **Inputs past the range** read the end code; the tool says so and the error then is the overshoot, not quantization. On the ADS1115 the  $\pm 4.096$  and  $\pm 6.144$  V settings are scaling only — the pins themselves never go past  $V_{\text{DD}} + 0.3$  V or below  $\text{GND} - 0.3$  V.
- **Single-ended on a bipolar part.** The ADS1115 measuring one input against ground uses only codes 0000h to 7FFFh, so it is a 15-bit converter in that mode. Model it on the bipolar pill with 16 bits, and read only the positive half.

## What it deliberately does not do

- **SNR, ENOB, oversampling.** The noise that quantization adds, and what averaging buys back, is the next tool, SNR estimation.
- **DAC.** The reverse conversion has its own tool.
- **Offset-binary and sign-magnitude codes.** Some bipolar converters output these instead of two's complement; they are rare enough on current parts to leave out.

## How it was checked

By hand, single-ended 12 bits on 3.3 V:  $\text{LSB} = 3.3 / 4096 = 805.66 \mu\text{V}$ ;  $1.2 \text{ V} / 805.66 \mu\text{V} = 1489.45$ , code 1489 = 0x5D1, code voltage 1.19963 V, error  $+366.2 \mu\text{V} = +0.4545$  LSB. Code 4095 gives 3.29919 V; 3.4 V reads 4095 and is flagged as over range.

Bipolar 16 bits on  $\pm 2.048$  V: LSB  $62.5 \mu\text{V}$ , matching the ADS1115's Table 7-1.  $-1.23456 \text{ V} / 62.5 \mu\text{V} = -19752.96$ , code  $-19753$ , register  $0xB2D7$  ( $65536 - 19753 = 45783$ ); error  $+2.5 \mu\text{V} = +0.04$  LSB. At  $\pm 4.096$  V the LSB is  $125 \mu\text{V}$ , again the datasheet's figure. Code  $-32768$  reads  $0x8000$  and  $-4.096$  V, the datasheet's  $-\text{FS}$ . 24 and 32 bits keep enough figures in the code voltage to tell neighbouring codes apart, and the hex fits its cell.


# DAC resolution

## What it computes

For an N-bit DAC, its reference and its output gain: the step between neighbouring output levels (the LSB), the code to write for a target voltage, the voltage that code really puts out, and how far that is from the target, in volts and in LSB. It works both ways — type the voltage you want, or type a code to see what it outputs. The output range is shown under the fields.

## Source

- **Microchip MCP4802/4812/4822** datasheet (DS20002249B):  
 $V_{OUT} = V_{REF} \times D / 4096$  with the GA bit at 1 ( $\times 1$ ), and  $2 \times V_{REF} \times D / 4096$  with GA at 0 ( $\times 2$ ), from the internal 2.048 V reference; its table gives 0.5 mV per step at  $\times 1$  and 1 mV at  $\times 2$ . The MCP4725 datasheet writes the same  $V_{REF} \times D_n / 4096$  with  $V_{REF} = V_{DD}$ .
- **ST's LL driver** for the STM32 DAC (`stm32f4xx_ll_dac.h`):  
`__LL_DAC_DIGITAL_SCALE` is 0xFFF, and  
`__LL_DAC_CALC_VOLTAGE_TO_DATA` multiplies by it and divides by  $V_{REF}$ , i.e.  $\text{code} = V \times 4095 / V_{REF}$ .



## Why the formulas are these

**One level per code.** A DAC puts out one voltage for each code and nothing in between:

$$V_{out} = \text{Code} \times \text{LSB}$$

so a target that is not exactly a level gets the nearer one, at most half an LSB away. That is why the picture is points, not a staircase: eight levels around the target at scale, the target a line between two of them, the chosen one in green.

**Two divisors, and why it matters.** How big the step is depends on what the datasheet divides by:

|                    |                                                       |                                                         |
|--------------------|-------------------------------------------------------|---------------------------------------------------------|
| $\div 2^N$ :       | $\text{LSB} = V_{REF} \times \text{Gain} / 2^N$       | top code = $V_{REF} \times \text{Gain} - 1 \text{ LSB}$ |
| $\div (2^N - 1)$ : | $\text{LSB} = V_{REF} \times \text{Gain} / (2^N - 1)$ | top code = $V_{REF} \times \text{Gain}$ exactly         |

Microchip, TI and Analog Devices write  $2^N$ : a resistor string or R-2R ladder of  $2^N$  equal steps, of which the top code uses  $2^N - 1$ , so the output never quite reaches VREF. ST's own code for the STM32 DAC uses  $2^N - 1$ , so 4095 is VREF itself. ST's documents do not all agree — a user on ST's community forum points out that RM0444 and AN3126 write 4096 — but the driver says 4095, and that user's bench measurements on a

NUCLEO-G071RB matched 4095 within a millivolt where 4096 was 3–5 mV out. The difference is under one LSB, but it is systematic, and it is always largest near the top of the range.

**Gain.** Many DACs follow the ladder with an amplifier. On the MCP4822 the GA bit picks  $\times 1$  or  $\times 2$  from its internal 2.048 V reference, which is how a 2.048 V reference makes a 4.095 V range at 1 mV a step.

## Assumptions and limits

- **The output cannot reach the rails.** The MCP4822's amplifier swings from 10 mV to  $V_{DD} - 40$  mV, and its accuracy holds only between those. The STM32 DAC with its output buffer on is similarly limited near 0 V and  $V_{REF+}$ . The top and bottom few codes on the tool's range are therefore not really available on those parts, and with  $\times 2$  gain the range can exceed  $V_{DD}$ , which the output will not follow either.
- **An ideal DAC.** Offset, gain error, INL and DNL are not modelled; they are specified in LSB, which is what this tool gives the size of.
- **$V_{REF}$  is taken as exact.** When it is  $V_{DD}$ , as on the MCP4725, the supply's tolerance scales every output.

## What it deliberately does not do

- **Bipolar DACs** ( $\pm V_{REF}$  output with offset-binary or two's-complement codes) are rare on current general-purpose parts; the bipolar pill on the ADC tool covers the arithmetic if needed.
- **PWM as a DAC.** Resolution there is set by the timer and the filter, not a ladder; it is a different calculation.
- **SNR and ENOB** are the next tool.

## How it was checked

$\div 2^N$ , 12 bits, 3.3 V: LSB 805.66  $\mu$ V; 1.2 V gives code 1489 (0x5D1), output 1.19963 V, error  $-366.2$   $\mu$ V =  $-0.4545$  LSB. Code 4095 outputs 3.29919 V, one LSB short of  $V_{REF}$ .

$\div (2^N - 1)$ : LSB  $3.3 / 4095 = 805.86$   $\mu$ V; code 4095 outputs exactly 3.3 V; 1.2 V gives code 1489, output 1.19993 V.

MCP4822 at  $\times 2$  ( $V_{REF}$  2.048 V, gain 2): LSB 1 mV, matching the datasheet's table; 3 V is code 3000 exactly; 5 V is flagged as above the 4.095 V range. 8 bits on the same settings: LSB 16 mV. 24 bits: output shown to enough figures to tell neighbouring codes apart.

# SNR estimation

## What it computes

Two things, one per pill.

- **Estimate** — the signal-to-noise ratio an ADC can reach from the noises that can be calculated: quantization, the process gain of oversampling, and sampling-clock jitter. Each is shown on its own, then added as noise powers into one SNR and expressed as an effective number of bits (ENOB).
- **ENOB ↔ SINAD** — the conversion between the two figures ADC datasheets quote, both ways, with how many of the converter's bits noise and distortion take away.

The picture is a level diagram in dBFS: full scale at the top, the signal, and each noise floor as a line. The gap from the signal down to the total floor is the SNR, and the highest noise line is the one worth fixing.

## Source

Walt Kester's Analog Devices tutorials:

- **MT-001**, "Taking the Mystery out of the Infamous Formula,  $\text{SNR} = 6.02N + 1.76 \text{ dB}$ ": the ideal SNR of an N-bit converter for a full-scale sine measured over dc to  $f_s/2$ , and the correction  $10 \log(f_s / 2BW)$  it calls process gain.
- **MT-003**, "Understand SINAD, ENOB, SNR, THD, THD + N, and SFDR": ENOB from SINAD, and its Equation 2, which adds the level below full scale so ENOB is normalised to full scale.
- **MT-007**, "Aperture Time, Aperture Jitter, Aperture Delay Time": the SNR limit set by jitter, and that the total jitter is the root-sum-square of the sampling clock's and the ADC's own aperture jitter.

The ENOB pill's defaults come from the **RP2040 datasheet**, section 4.9.3: SINAD 54.0 dB typical at 997 Hz and almost full scale, ENOB 8.7.


## Why the formulas are these

**Quantization.** Rounding to the nearest code leaves an error spread evenly over  $\pm \frac{1}{2}$  LSB, whose rms value is  $q / \sqrt{12}$ . A full-scale sine has an rms value of  $2^N q / (2\sqrt{2})$ . Their ratio, in dB, is

$$\text{SNR} = 6.02N + 1.76 \text{ dB}$$

— 6.02 dB per bit, which is  $20 \log 2$ , plus a constant from the shape of a sine. 12 bits: 74.0 dB. 16 bits: 98.1 dB.

**Level.** The quantization noise stays where it is when the signal gets smaller, so the SNR falls one dB for every dB below full scale. A 12-bit ADC reading a signal at  $-20$  dBFS has 54 dB of SNR, the same as an ideal 8.7-bit



one. That is the most common way resolution is wasted.

**Process gain.** Quantization noise is spread evenly from dc to  $f_s/2$ . Filter the result digitally down to a bandwidth BW and only the share of the noise inside BW remains:

$$PG = 10 \log(f_s / 2BW)$$

Four times the bandwidth needed is 6.02 dB, one extra bit. MT-001's example: a 65 MSPS ADC with channels 30 kHz wide gains 30.3 dB, 65 dB becoming 95.3 dB.

**Jitter.** If the instant of sampling wanders by  $t_j$  rms, the error it causes is the signal's slope times that wander. For a full-scale sine at  $f_{in}$  the result is independent of amplitude:

$$SNR(jitter) = -20 \log(2\pi \times f_{in} \times t_j)$$

It depends on the input frequency, not the sample rate. 100 ps is harmless at 1 kHz (124 dB) and ruinous at 100 MHz (24 dB). MT-007 notes that 14-bit performance at 100 MHz needs under 0.1 ps; the tool gives 84 dB at exactly 0.1 ps, just under the 86 dB of 14 bits.

**Adding them.** Uncorrelated noises add as powers:

$$SNR = -10 \log(10^{(-Q/10)} + 10^{(-J/10)})$$

so the total is always a little worse than the worse of the two, and a term 10 dB better than the other barely counts. When jitter is the lower one, the tool says so — more bits will not help.

**What ENOB means.** ENOB, the *effective number of bits*, answers the question "how many bits is this converter really worth?". The resolution on the datasheet, 12 bits say, only counts the codes the ADC can output. A real ADC also adds its own noise and distortion, so the last few bits are buried in them: they change from reading to reading without telling you anything about the signal. ENOB is the resolution of a perfect converter that would be exactly as noisy as the real one. A 12-bit ADC with an ENOB of 8.7 gives, on a changing signal, the same quality as a flawless 8.7-bit ADC; the bottom 3.3 bits are noise. ENOB is not a whole number, because it is worked back from a measured ratio, not counted.

**What SINAD means.** SINAD, the *signal-to-noise-and-distortion ratio*, is that measured ratio: a pure sine is fed in, and the power of the sine is compared with the power of everything else in the output — noise and the harmonics that distortion adds. SNR is the same measurement leaving the harmonics out, which is why a datasheet's SNR is always the higher of the two. Because a perfect N-bit converter has an SINAD of  $6.02N + 1.76$  dB, the two figures are the same fact in different units, and datasheets quote either.

**ENOB from the ratio.** The SNR equation solved for N, with the level term of MT-003:

$$ENOB = (SNR - 1.76 - Level) / 6.02$$

On the ENOB pill SINAD takes the place of SNR, as the datasheets do. The RP2040's 54.0 dB gives 8.68 bits, and its datasheet says 8.7: of its 12 bits, 3.3 are lost to noise and distortion. Its SNR is 61.5 dB but its THD is

–55 dB — its INL and DNL errors (the DNL spikes are erratum RP2040-E11) cost more than the noise.

## Assumptions and limits

- **Oversampling needs noise.** Process gain assumes the quantization error is random. A clean, slow signal on a quiet ADC can give the same code every time, and averaging a constant gives nothing. MT-001 warns about this correlation; the cure is noise or dither of about an LSB at the input.
- **Thermal noise and distortion are not estimated.** They are the ADC's own and only its datasheet knows them. The Estimate pill gives the ceiling the architecture allows; the ENOB pill gives what the part actually does.
- **No noise shaping.** Delta-sigma converters push quantization noise out of band and gain far more than 10 log per octave of oversampling. Their datasheets give SNR directly.
- **SINAD at other levels.** MT-003's level correction assumes the noise does not change with level. Distortion usually drops at lower levels, so ENOB normalised from a –6 dBFS measurement can read better than at full scale.

## What it deliberately does not do

- **THD, SFDR, THD + N.** Read from the datasheet, not calculated.
- **Noise-free resolution and effective resolution** from rms input noise, the measure for dc and slow signals. MT-003 warns not to confuse them with ENOB; they are a different calculation from different data.

## How it was checked

Default Estimate: 12 bits, 1 MSPS, 500 kHz, 0 dBFS, 100 kHz, 100 ps.  $Q = 6.02 \times 12 + 1.76 = 74.0$  dB;  $J = -20 \log(2\pi \times 10^5 \times 10^{-10}) = 84.04$  dB; total 73.59 dB; ENOB 11.93 bits.

MT-001's process-gain example: 65 MSPS and 30 kHz give 30.35 dB. MT-007's jitter example: 14 bits, 100 MHz, 0.1 ps give  $J = 84.04$  dB, total 81.91 dB, with the jitter warning. 40 MHz of bandwidth at 65 MSPS is refused as past  $f_s/2$ . With no jitter at –1 dBFS, 14 bits give 85.04 dB and an ENOB of 14.0, normalised.

ENOB pill: SINAD 54.0 dB at 0 dBFS gives 8.678 bits — the RP2040's 8.7 — with 3.32 bits lost; ENOB 10 gives SINAD 61.96 dB.